

NUMERICAL ANALYSIS FOR ARTIFICIAL INTELLIGENCE, WEEK 4

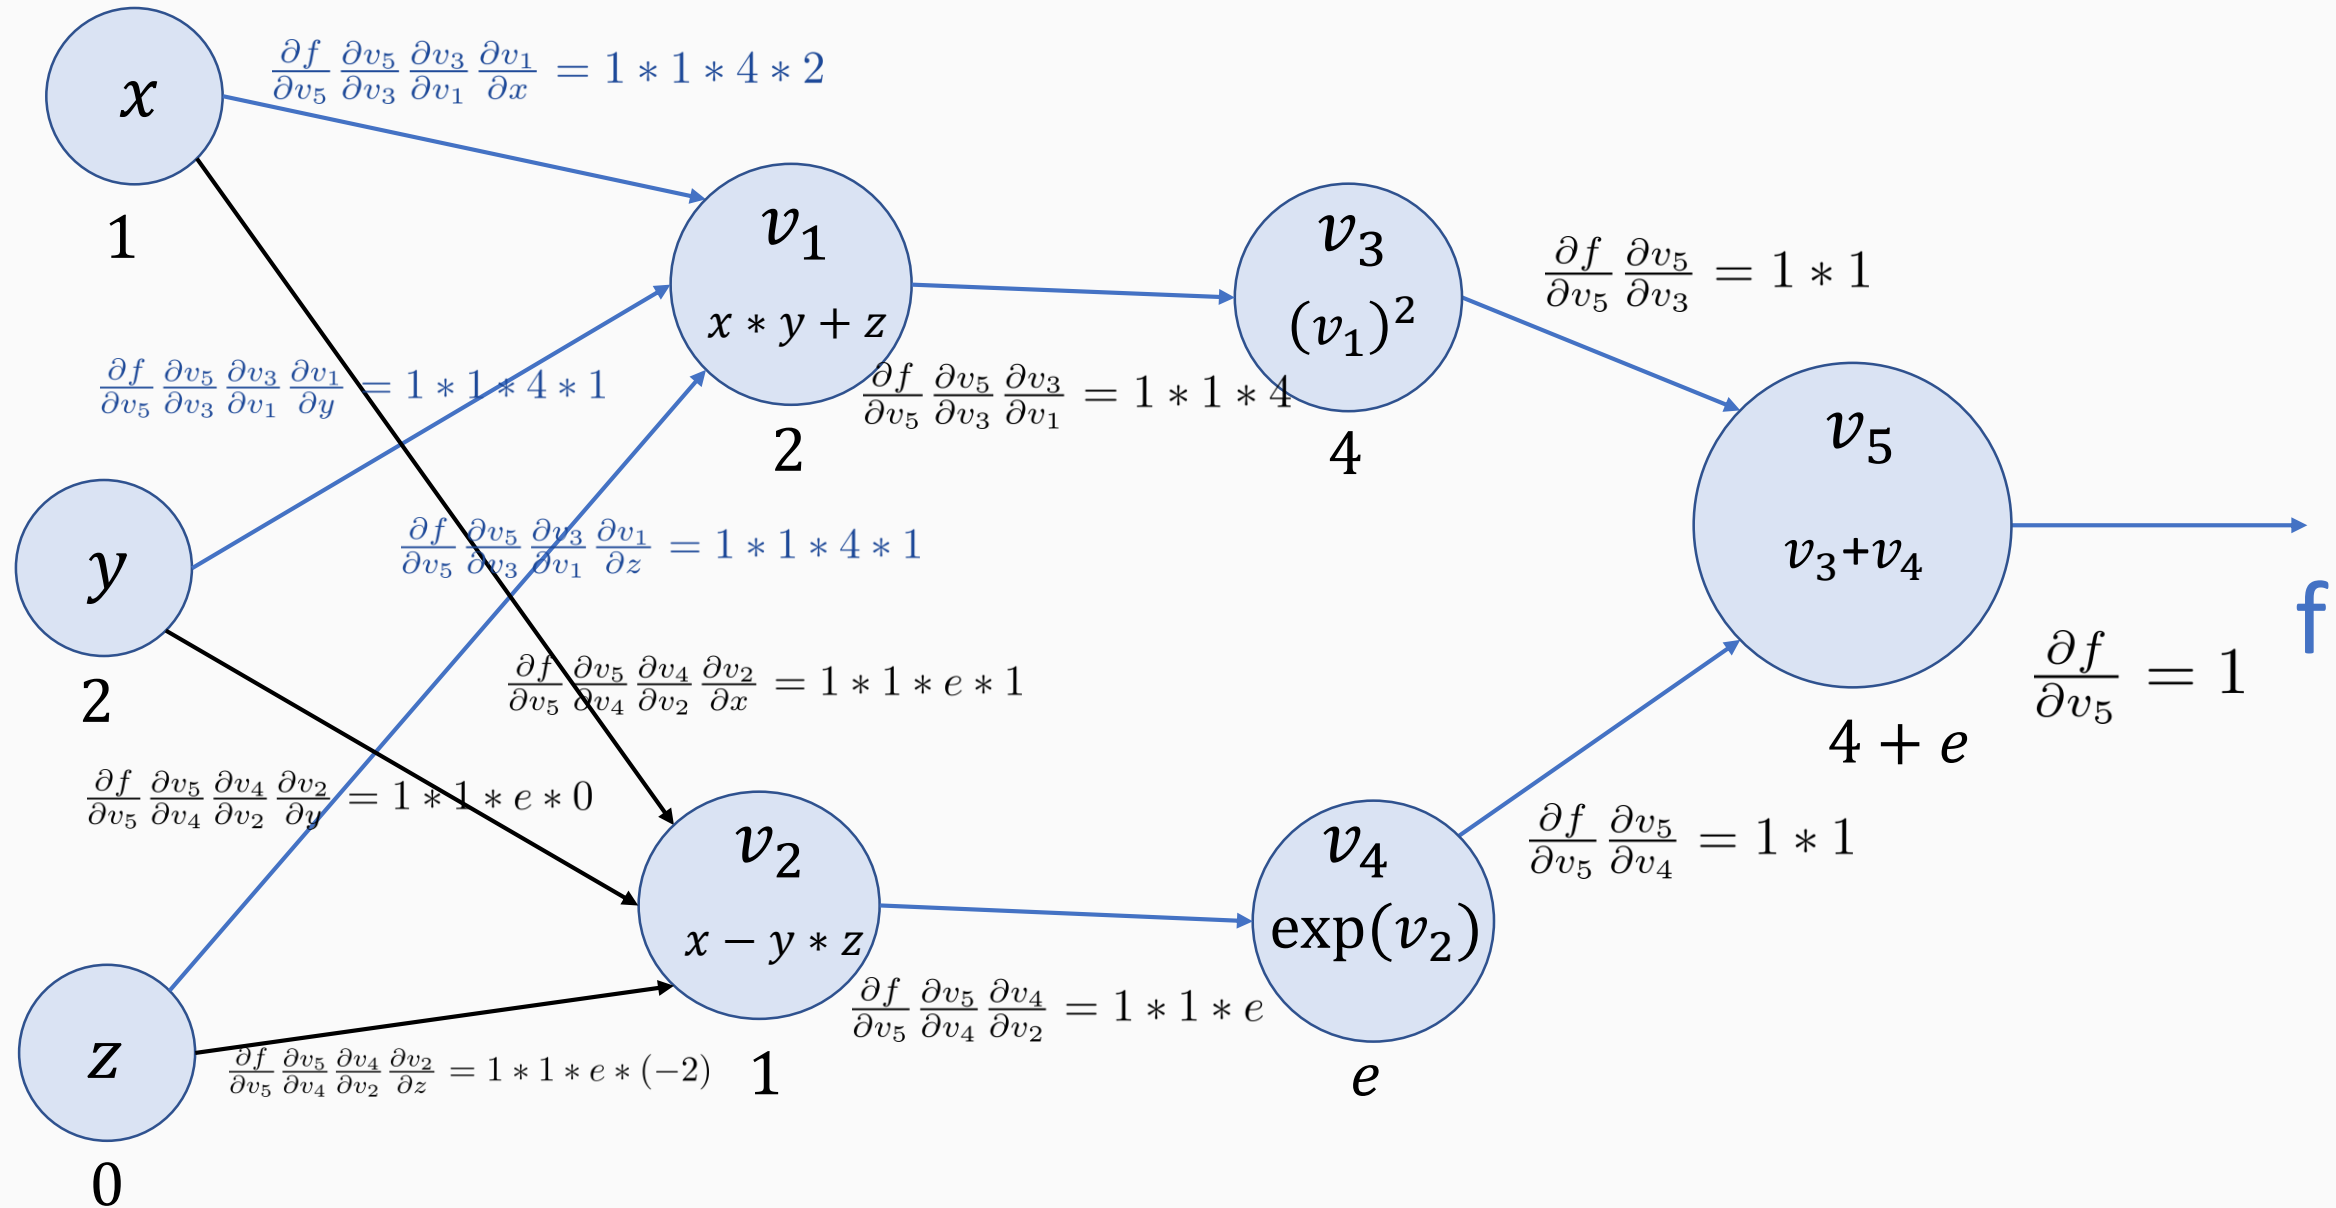
UCSD Summer session II 2018

CSE 190

Jacek (*Yatzek*) Cyranka

Backprop of $f(x,y)=(xy + z)^2 + \exp(x - yz)$

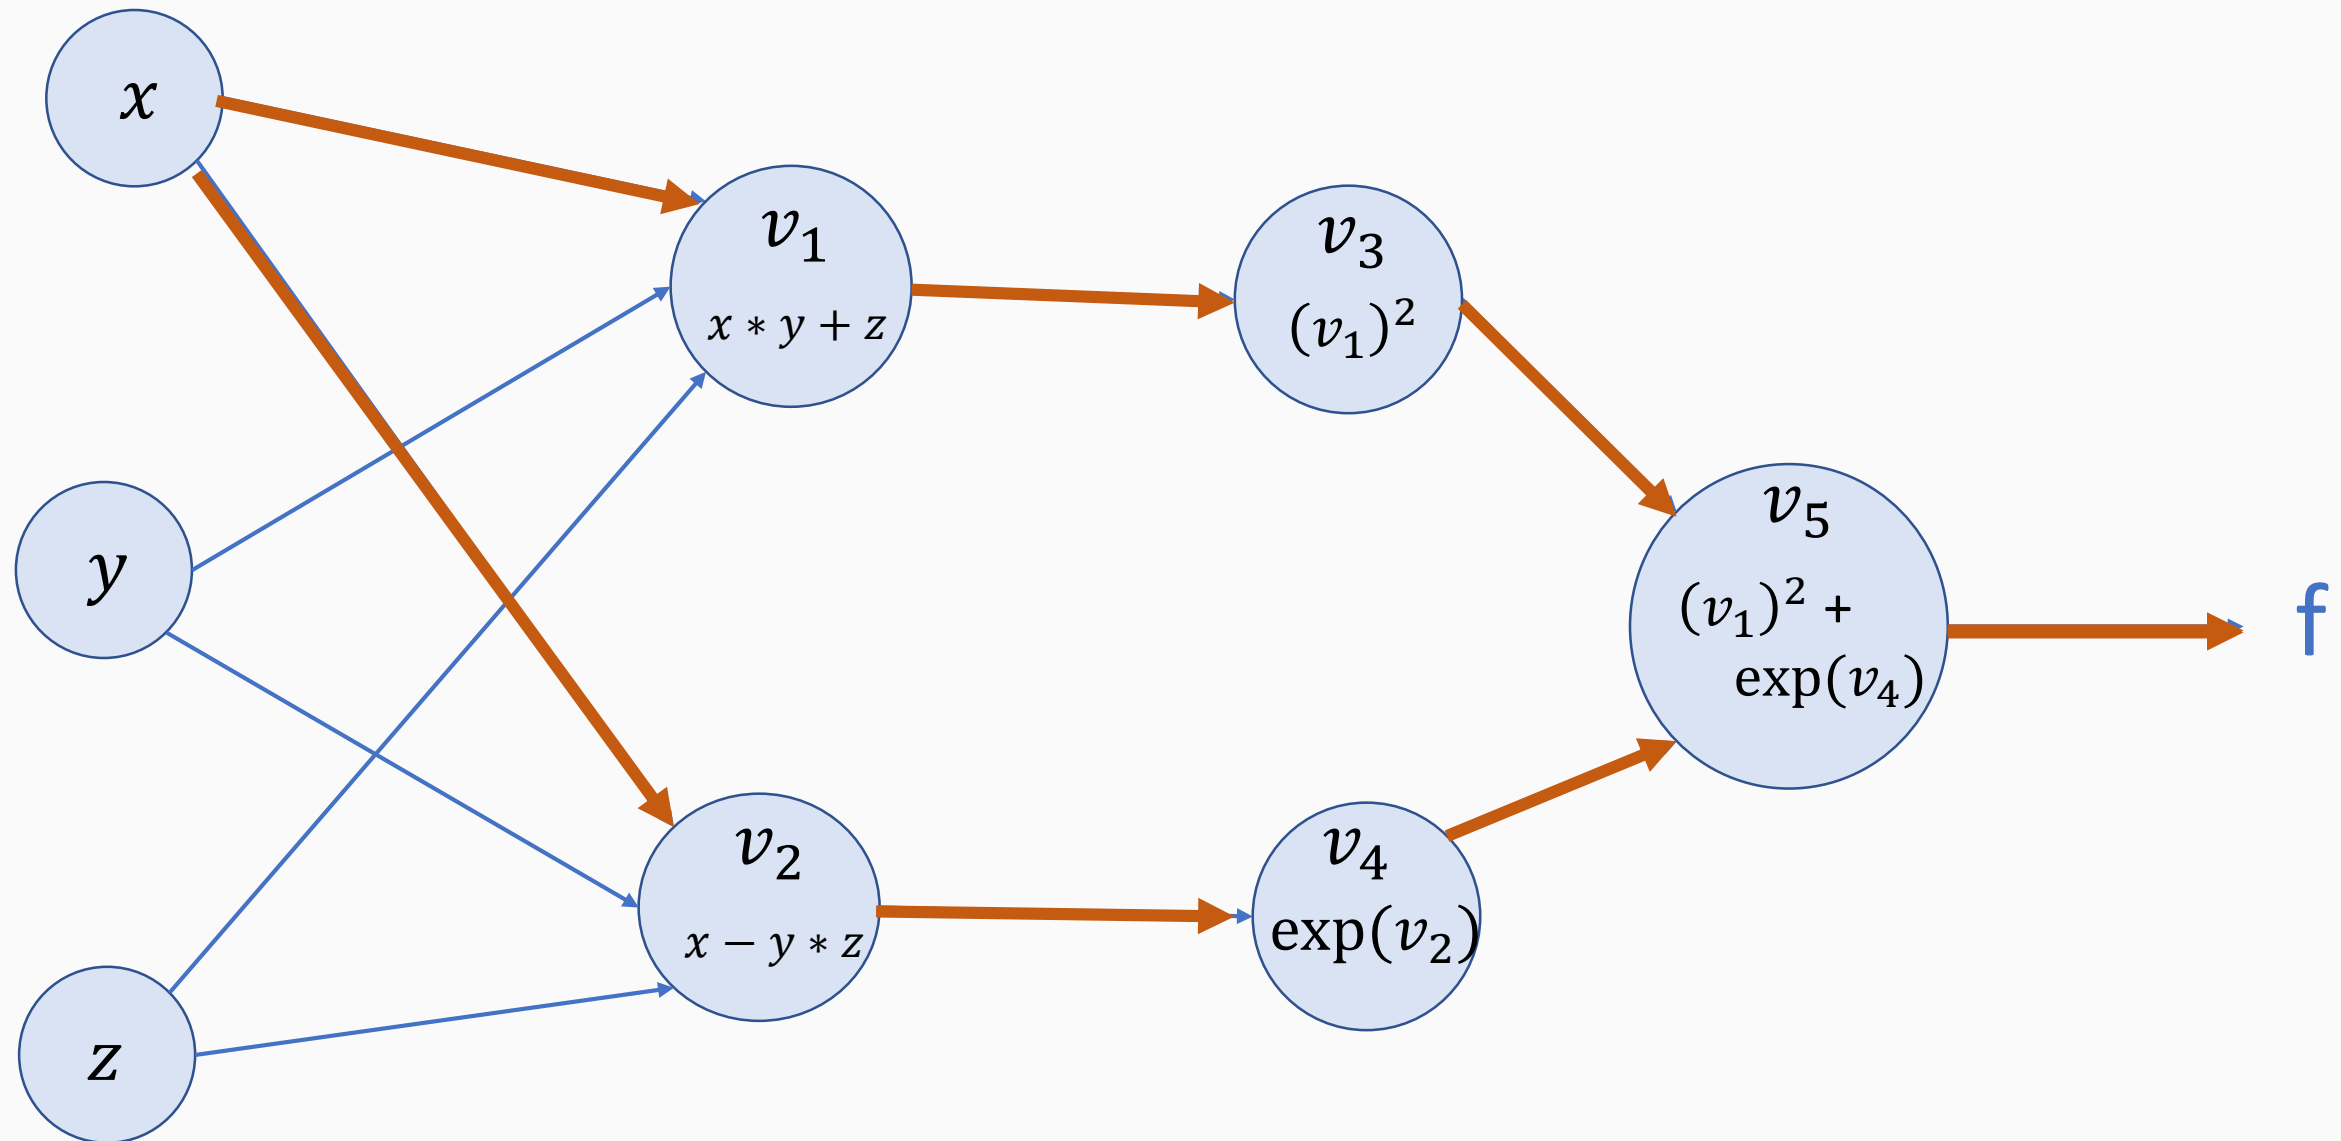
A more complicated graph (function)



Interpretation of backprop as sum of paths

$$\frac{\partial f}{\partial x} = \sum_{\substack{\text{path in the graph} \\ \text{from } x \text{ to } f \\ (v_{i_1}, v_{i_2}, \dots, v_{i_k})}} \frac{\partial f}{\partial v_{i_k}} \frac{\partial v_{i_k}}{\partial v_{i_{k-1}}} \dots \frac{\partial v_{i_2}}{\partial v_{i_1}}$$

Partial derivatives – paths interpretation



Hence the total partial derivatives are

We combine the results obtained by ‘traversing’ all paths in the graph, and the final result for $f(x,y)=(xy + z)^2 + \exp(x - yz)$ is


$$\frac{\partial f}{\partial x} = \frac{\partial f}{\partial v_5} \frac{\partial v_5}{\partial v_3} \frac{\partial v_3}{\partial v_1} \frac{\partial v_1}{\partial x} + \frac{\partial f}{\partial v_5} \frac{\partial v_5}{\partial v_4} \frac{\partial v_4}{\partial v_2} \frac{\partial v_2}{\partial x} = 1 * 1 * 4 * 1 + 1 * 1 * e * 1,$$

$$\frac{\partial f}{\partial y} = \frac{\partial f}{\partial v_5} \frac{\partial v_5}{\partial v_3} \frac{\partial v_3}{\partial v_1} \frac{\partial v_1}{\partial y} + \frac{\partial f}{\partial v_5} \frac{\partial v_5}{\partial v_4} \frac{\partial v_4}{\partial v_2} \frac{\partial v_2}{\partial y} = 0 + 1 * 1 * 4 * 2,$$

$$\frac{\partial f}{\partial z} = \frac{\partial f}{\partial v_5} \frac{\partial v_5}{\partial v_4} \frac{\partial v_4}{\partial v_1} \frac{\partial v_1}{\partial z} + \frac{\partial f}{\partial v_5} \frac{\partial v_5}{\partial v_4} \frac{\partial v_4}{\partial v_2} \frac{\partial v_2}{\partial z} = 1 * 1 * 4 * 1 + 1 * 1 * e * (-2).$$

Multivariate sigmoid

$$\text{Backprop of } f = \frac{1}{1 + e^{-(w_0 x_0 + w_1 x_1 + w_2 x_2)}}$$



Backpropagation using vectorization

If nodes have vector values, use the generalized chain rule

$$\frac{\partial z}{\partial x_k} = \sum_{j=1}^n \frac{\partial z}{\partial y_j} \frac{\partial y_j}{\partial x_k},$$

$$\nabla_x z = \left(\frac{\partial y}{\partial x} \right)^T \nabla_y z. \quad \leftarrow \text{Gradient (column vec.)}$$


z is a scalar, $x \in \mathbb{R}^m$, $y \in \mathbb{R}^n$, then

$$\frac{\partial y}{\partial x} \in \mathbb{R}^{n \times m} \quad \leftarrow \text{Jacobian matrix}$$

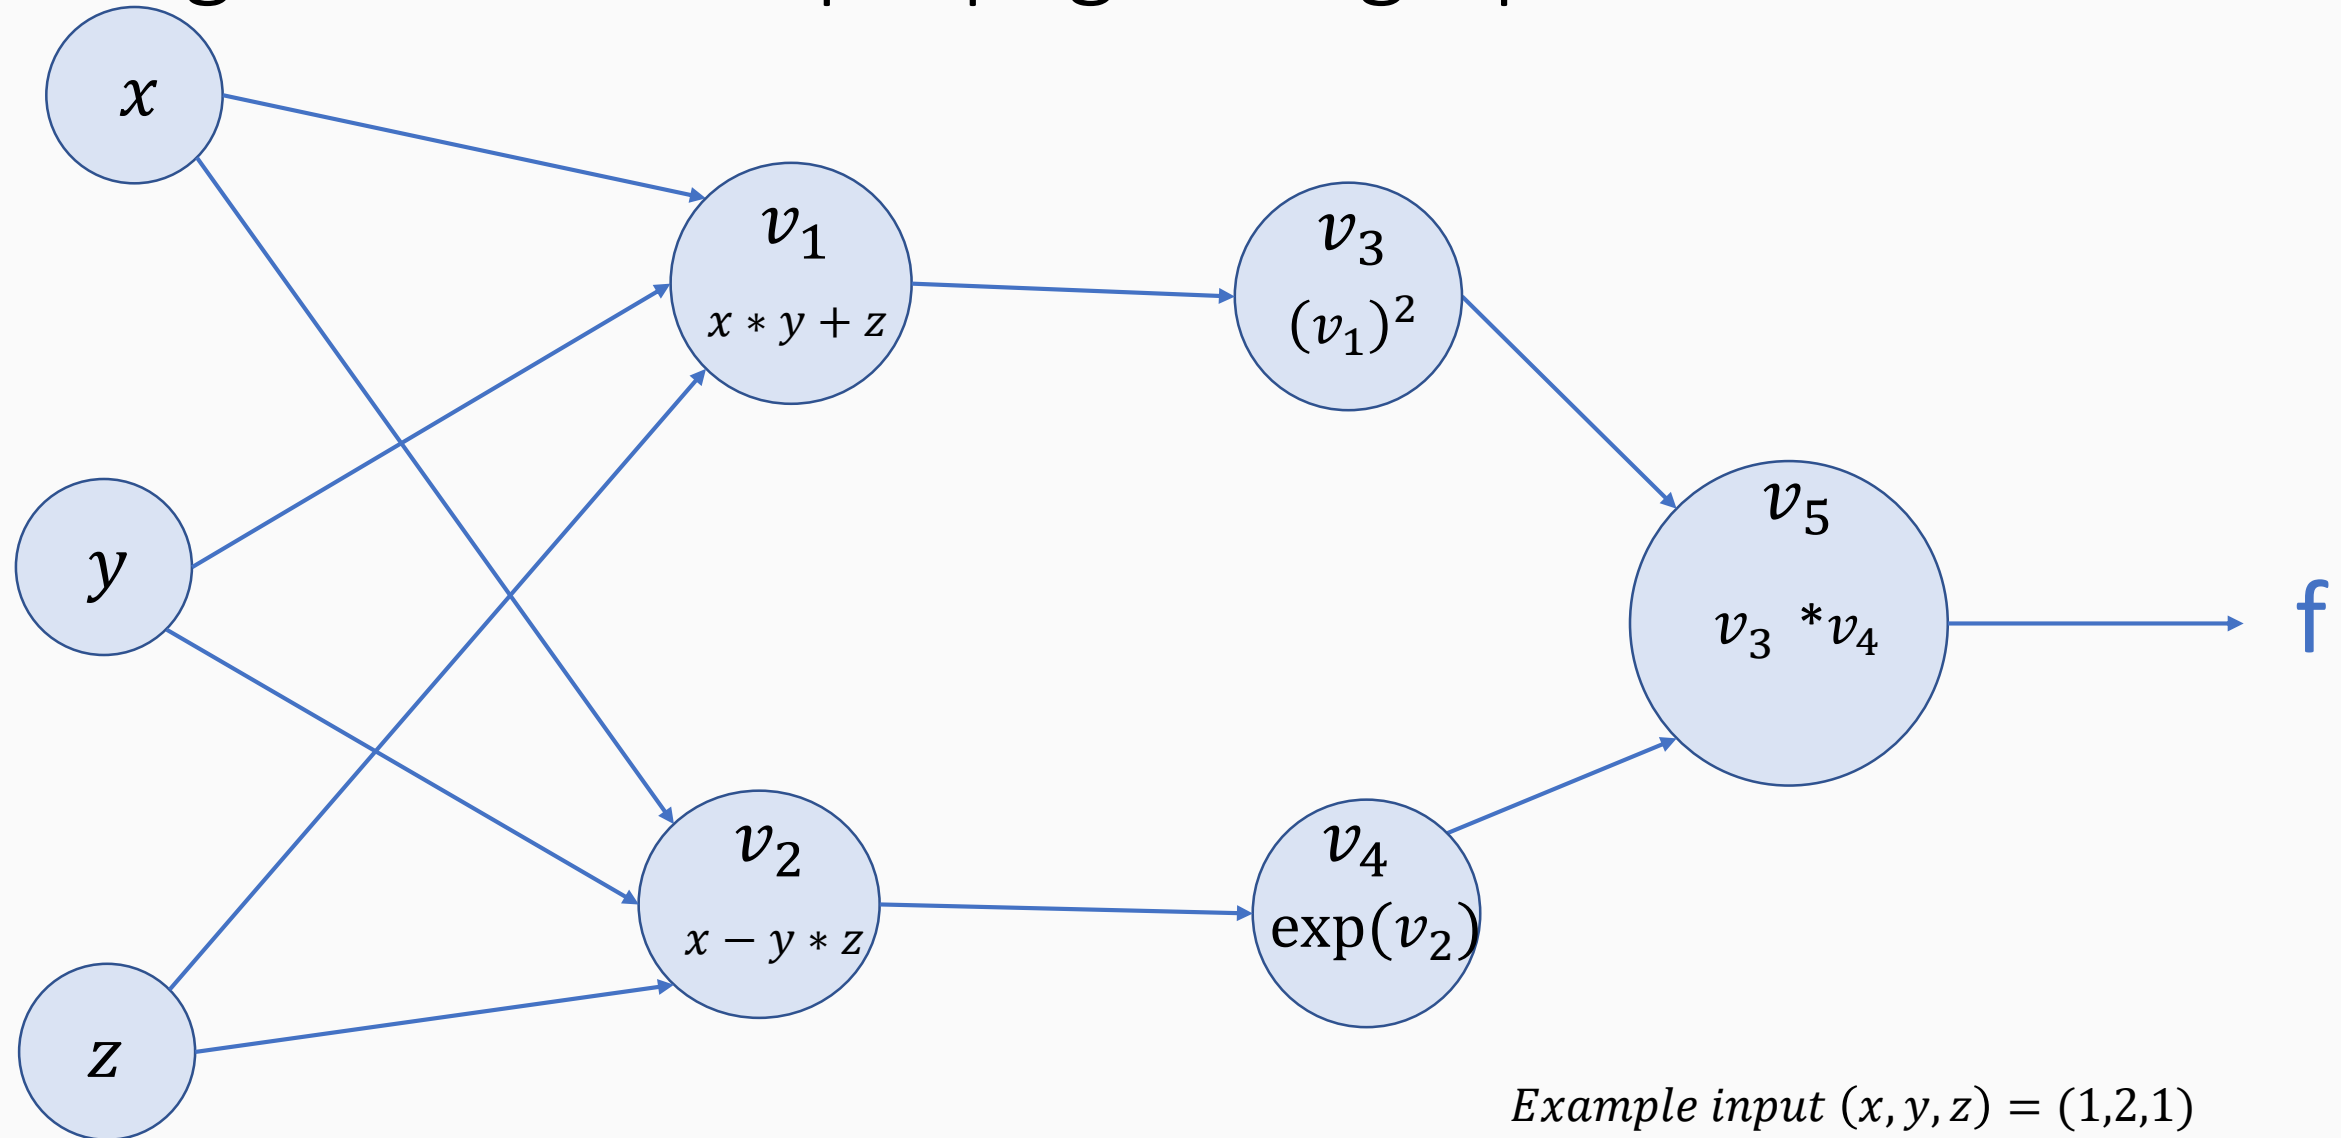
Vectorized backprop case study

$$\text{Backprop of } f = \frac{1}{1+e^{-(x^T w + b)}} = \sigma(x^T w + b)$$

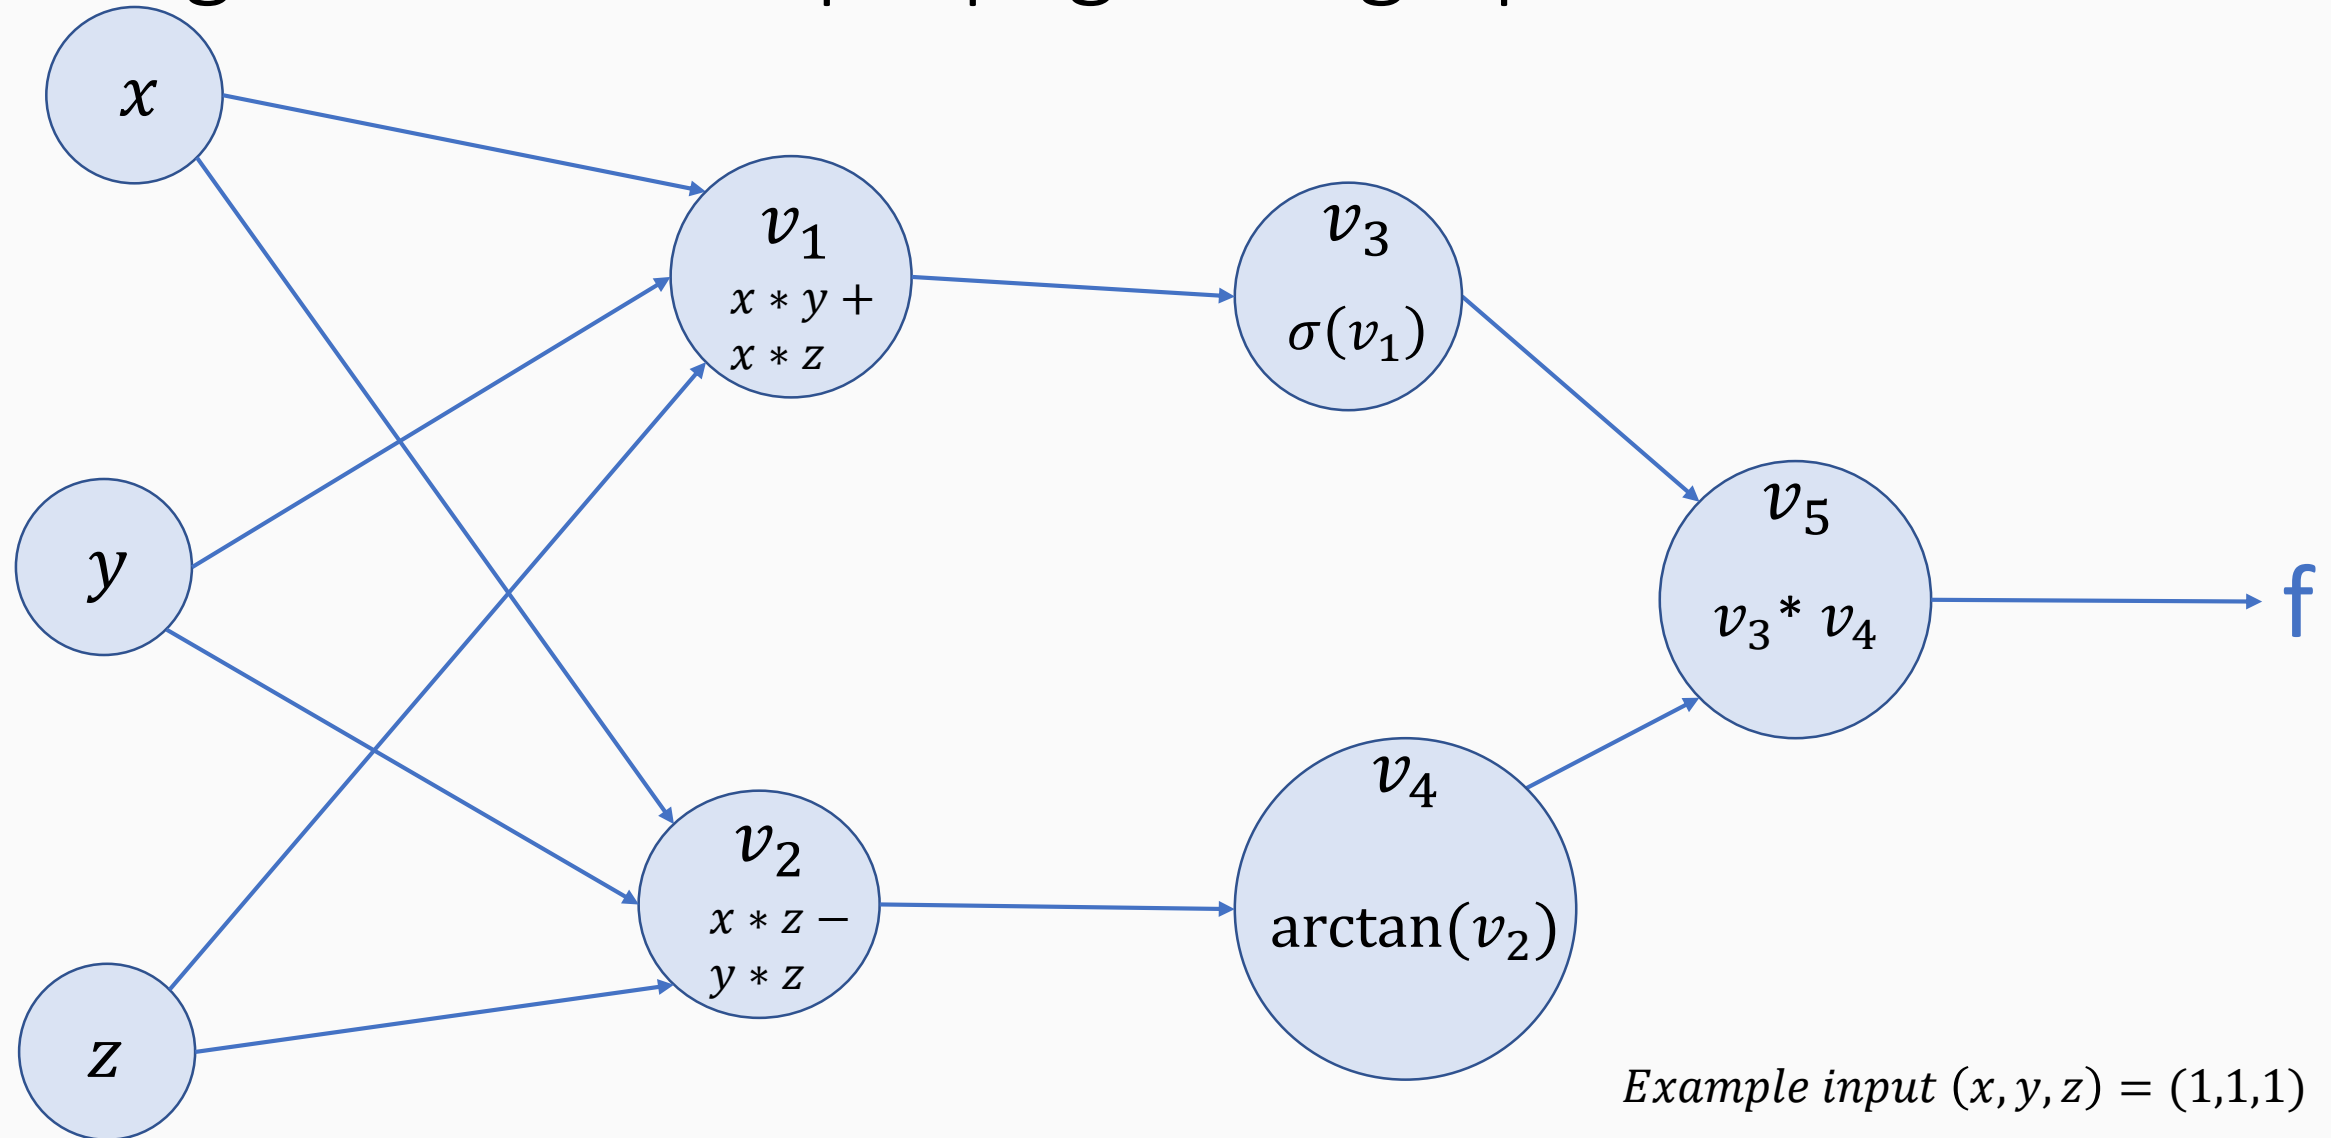
$x, w \in \mathbb{R}^n$ are vectors, b is a number



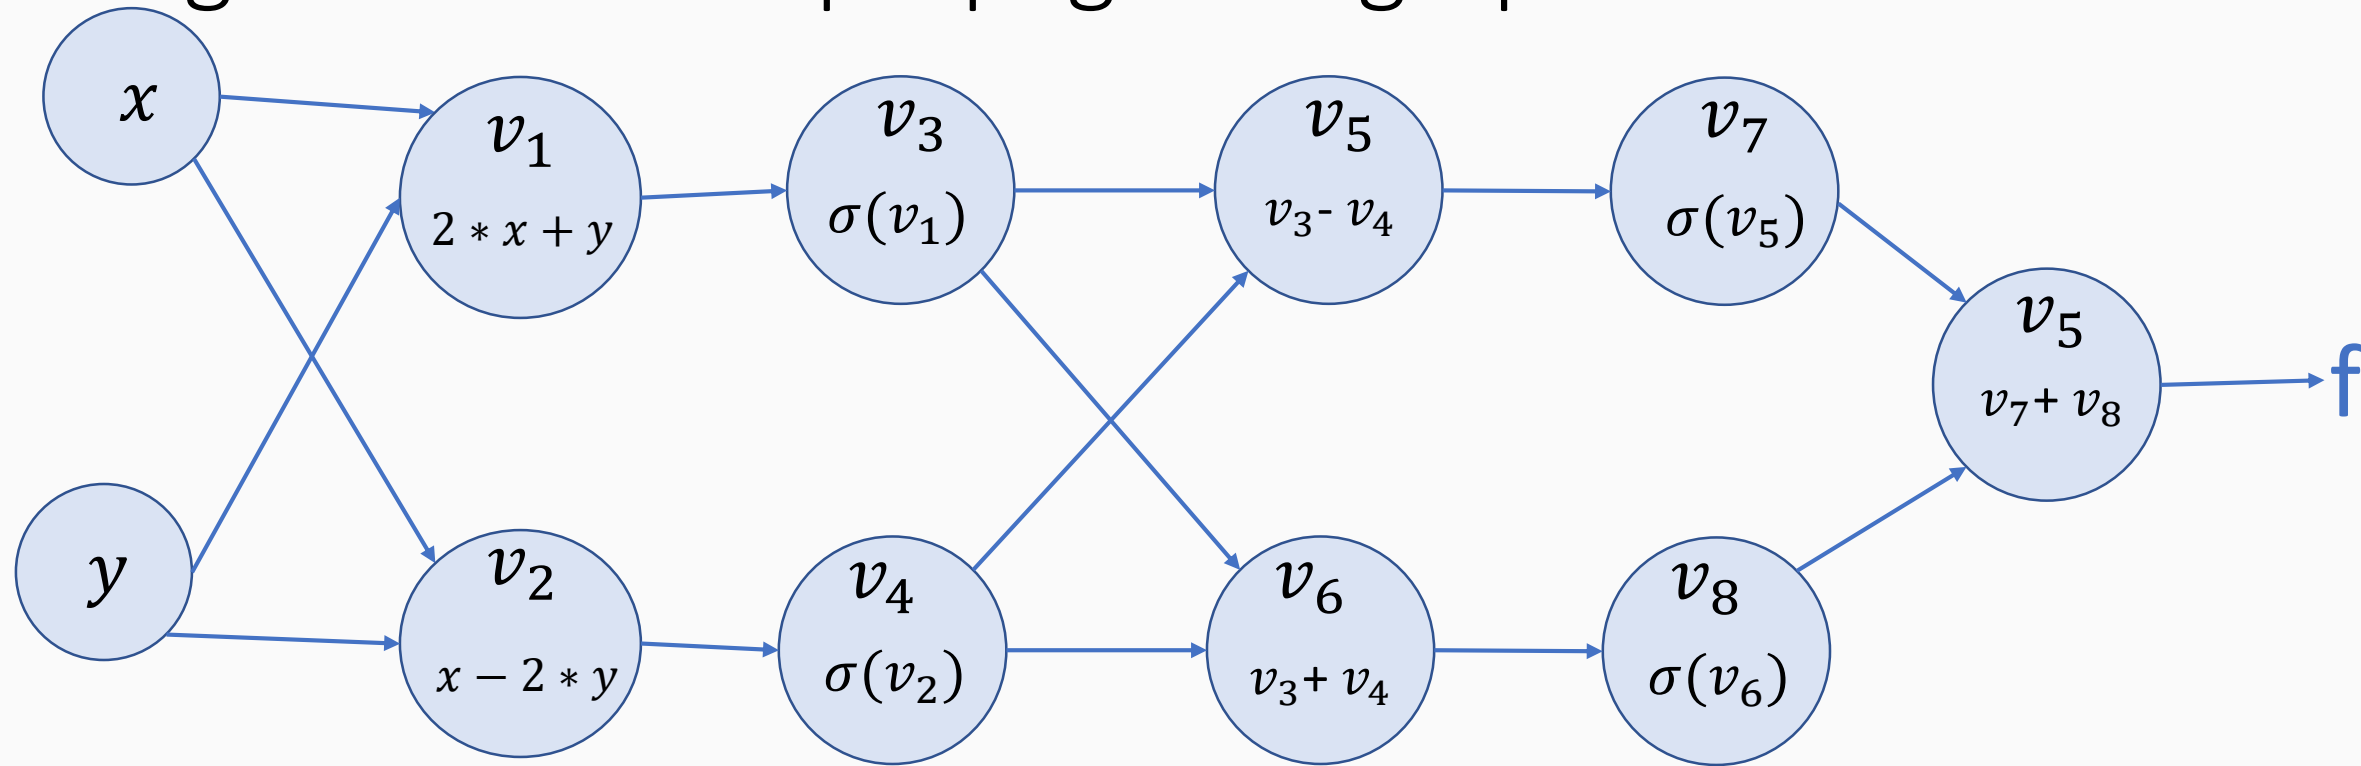
Assignment 6 backpropagation graph 1



Assignment 6 backpropagation graph 2

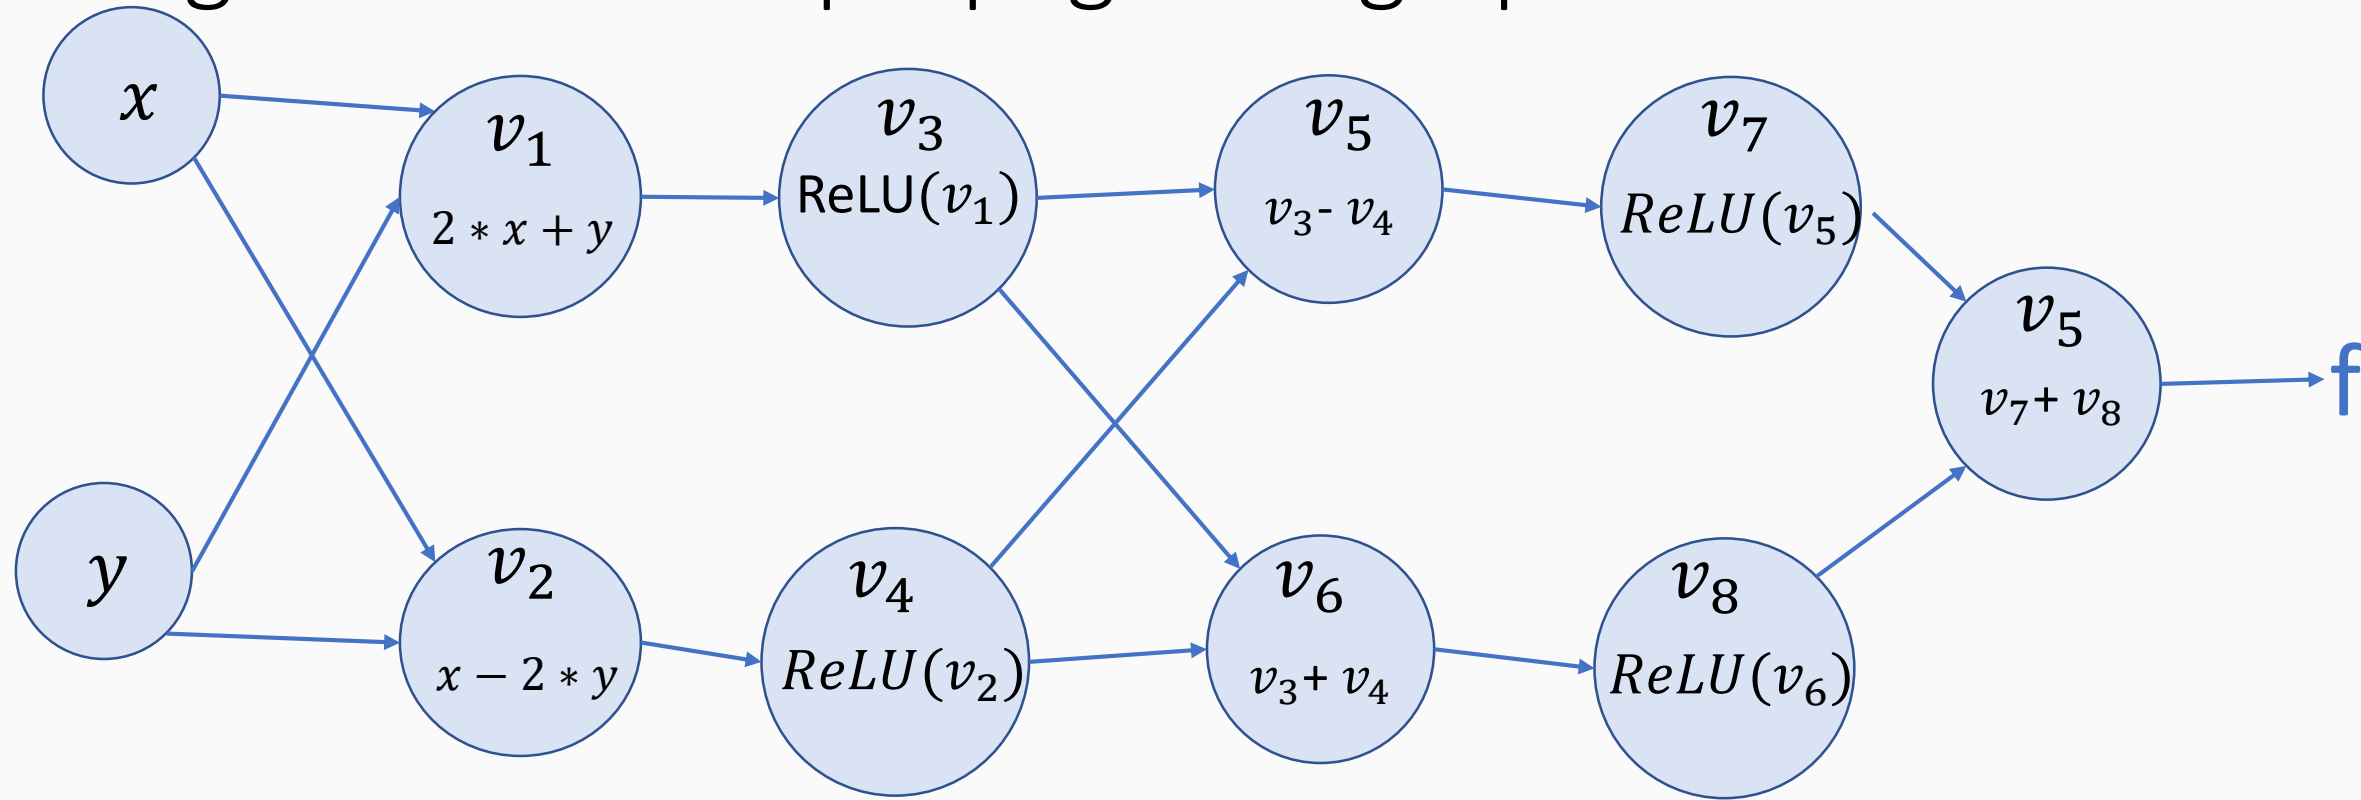


Assignment 6 backpropagation graph 3



Example input $(x, y) = (1, 1)$

Assignment 6 backpropagation graph 4



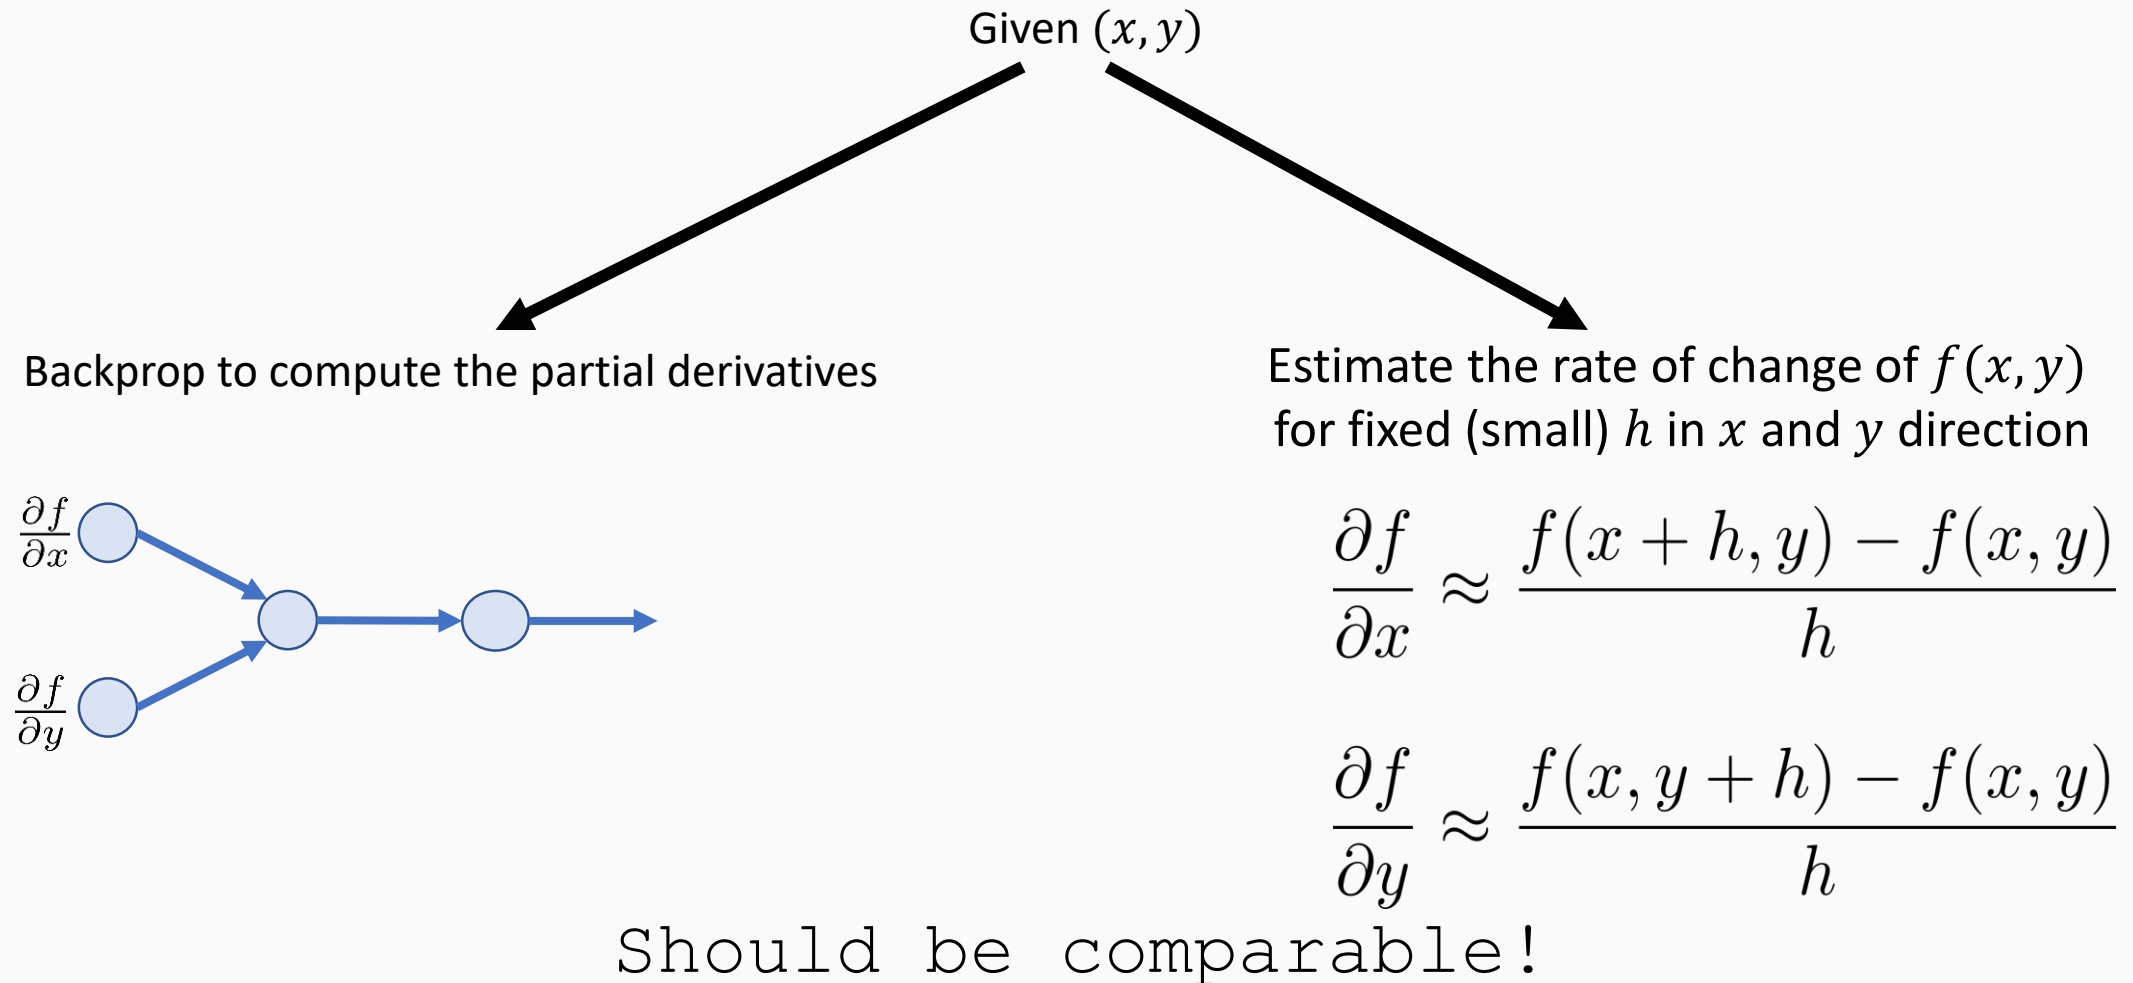
Example input $(x, y) = (1, 1)$

$$\text{ReLU}(x) = \begin{cases} x & \text{if } x > 0, \\ 0 & \text{if } x \leq 0 \end{cases},$$



$$\text{ReLU}'(x) = \begin{cases} 1 & \text{if } x > 0, \\ 0 & \text{if } x \leq 0 \end{cases}$$

Backpropagation gradient checking


How to debug my backprop algorithm implementation ???



Biological vs artificial neural nets



Original perceptron model



Original perceptron used the *Heaviside activation function*

$$f(x) = \begin{cases} 1, & \text{if } w \cdot x + b > 0, \\ 0, & \text{otherwise} \end{cases}$$

Rosenblatt, Frank (1957), *The Perceptron--a perceiving and recognizing automaton*. Report 85-460-1, Cornell Aeronautical Laboratory.

Separates two regions of the space using linear (hyper)plane, can perform linear separable tasks.

1	0	1
0	0	0
<i>AND</i>	0	1

1	1	1
0	0	1
<i>OR</i>	0	1

1	1	0
0	1	0
<i>NOT</i>	0	1


However, not all logic operations are linearly separable

1	1	0
0	0	1
<i>XOR</i>	0	1

Need to use two hidden layers to solve XOR


Why training NN is hard ???

The function of NN weights that is being optimized is nonconvex



Hardness of nonconvex optimization

1. There can be local minima, which are way worse than global,
2. There can be saddle-points,
3. No control where gradient descent will converge to,
4. Gradient descent converge to different critical points depends on initial points.



Why training NN is hard ??? (NP-complete)

Training a 3-Node Neural Network is NP-Complete


*Avrim L. Blum and Ronald L. Rivest**

MIT Laboratory for Computer Science
Cambridge, Massachusetts 02139

`avrim@theory.lcs.mit.edu`


`rivest@theory.lcs.mit.edu`

1. NP-complete problem means any provided example solution can be verified in polynomial time.
2. For any NP problem the best known algorithm of finding a solution runs in polynomial time.
3. Any problem in NP can be reduced to a NP-complete problem in polynomial time.
4. If there exists a polynomial time algorithm for solving any of NP-complete problems, then $P=NP$.



NP-completeness cont.

Training of the three node neural network to perform AND on the outputs of N_1, N_2 is NP-complete.



Where each node N_i computes a linear threshold function

$$N_i(x) = \begin{cases} +1 & \text{if } a_1x_1 + a_2x_2 + \cdots + a_mx_m > a_0, \\ -1 & \text{otherwise} \end{cases}$$


Show that this problem is NP-complete by reduction to the problem of *Set-Splitting* known to be NP-complete

The following problem, *Set-Splitting*, was proven to be NP-complete by Lovász (see Garey and Johnson [6]).

“Given a finite set S and a collection C of subsets c_i of S , do there exist disjoint sets S_1, S_2 such that $S_1 \cup S_2 = S$ and for each i , $c_i \not\subseteq S_1$ and $c_i \not\subseteq S_2$?”

NP-completeness cont.

1. Arrange all training examples on a n -dimensional hypercube



2. Label all nodes (training examples) $+/-$ according to the set-splitting instance

- Let the origin 0^n be labeled '+'.
– For each s_i , put a point labeled '-' at the neighbor to the origin that has a 1 in the i th bit: that is, at $(00 \cdots 010 \cdots 0)$. Call this point \mathbf{p}_i .
– For each $c_j = \{s_{j1}, \dots, s_{jk_j}\}$, put a point labeled '+' at the location whose bits are 1 at exactly the positions j_1, j_2, \dots, j_{k_j} : that is, at $\mathbf{p}_{j1} + \dots + \mathbf{p}_{jk_j}$.

Set-splitting example corresponding to the cube above is

$$S = \{s_1, s_2, s_3\}, c_1 = \{s_1, s_2\}, c_2 = \{s_2, s_3\}.$$

Hence, '-' for nodes 001, 010, 100, and '+' for nodes 000, 011, 110.

3. There is a solution of set-splitting problem \leftrightarrow there exists two hyperplanes such that all positive/negative nodes are separated within one of the quadrants.

Sketch the NP-completeness proof

one way \Rightarrow

1. Given S_1, S_2 from the solution of to the Set-Splitting. Define planes
 $P_1: a_1x_1 + \dots + a_nx_n = -\frac{1}{2}$, where $a_i = -1$ if $s_i \in S_1$, and $a_i = n$ if $s_i \notin S_1$
And similarly
 $P_2: b_1x_1 + \dots + b_nx_n = -\frac{1}{2}$, where $b_i = -1$ if $s_i \in S_2$, and $b_i = n$ if $s_i \notin S_2$
2. P_1 separates all $-$ nodes from the origin, because they evaluate to $-1 < -\frac{1}{2}$, and also P_2 separates all $-$ nodes from the origin as they evaluate to $-1 < -\frac{1}{2}$ also.
3. Hence, the quadrant $P_1 > -\frac{1}{2}, P_2 > -\frac{1}{2}$ contains $+$ nodes exclusively and therefore the three-node network can be trained to perform $\text{AND}(N_1, N_2)$ operation.

Sketch of the NP-completeness proof II

the other way \Leftarrow

1. Given two planes P_1, P_2 , let S_1 be the set of points separated from the origin by P_1 ($-$ nodes), and S_2 be the set of points sep. from the origin by P_2 ($-$ nodes).
2. It holds $S = S_1 \cup S_2$, as all $-$ nodes are separated by either P_1 or P_2 .
3. Consider $c_j = \{s_{j_1}, \dots, s_{j_k}\}$, if say $c_j \subset S_1$, then P_1 must separate all $p_i = (0 \dots 0 \overset{\uparrow j_i}{1} 0 \dots 0)$ from the origin.
4. But then P_1 must separate also $p_1 + p_2 + \dots + p_k$ ($+$ node) which means that $+$ nodes are not confined to a single quadrant, which contradicts the assumption.

TOPIC:

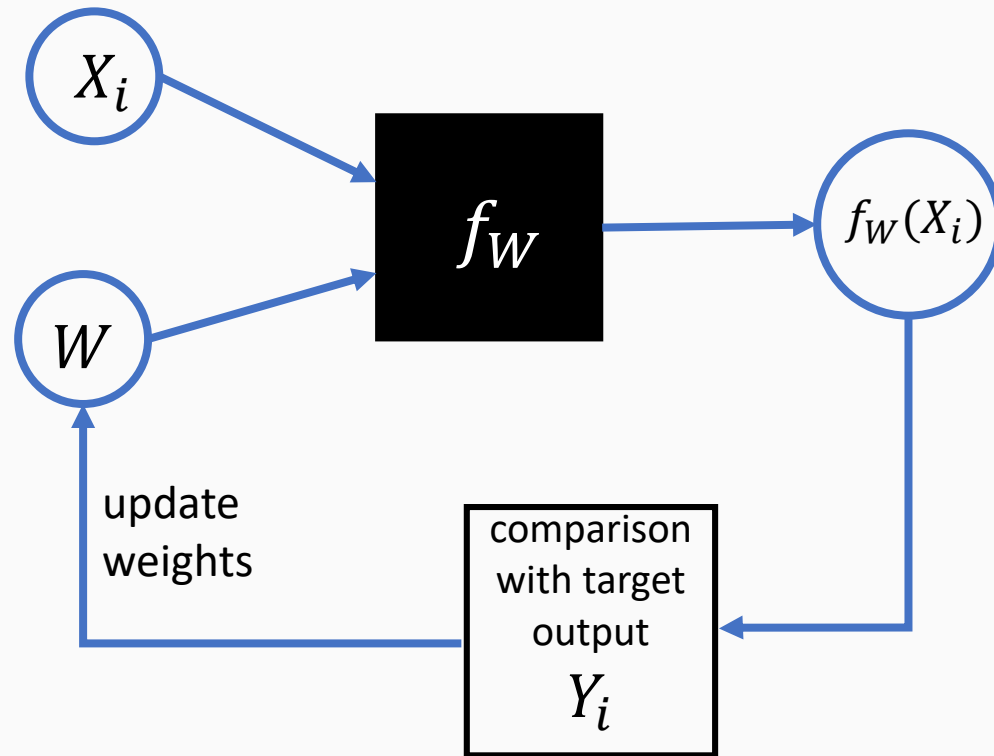
SINGLE HIDDEN LAYER NEURAL
NETWORK, SUPERVISED LEARNING

General algorithm of neural network supervised learning

In supervised learning there need to be a teacher

Teacher in the context of NN \equiv set of N training pairs $\{(X_i, Y_i)\}_{i=1}^N$

For each i , and training pair (X_i, Y_i)



Goal is to adjust W , such that the loss value

$$L(W) = \frac{1}{N} \sum_{i=1}^N \|f_W(X_i) - Y_i\|_2^2 \text{ is minimized.}$$

New convention of indexing

As we start dealing with multiple layers and tensors we need to introduce a new convention for indexing

position in sequence
(layer number)

W^i weight matrix,

indices (row and
column #)

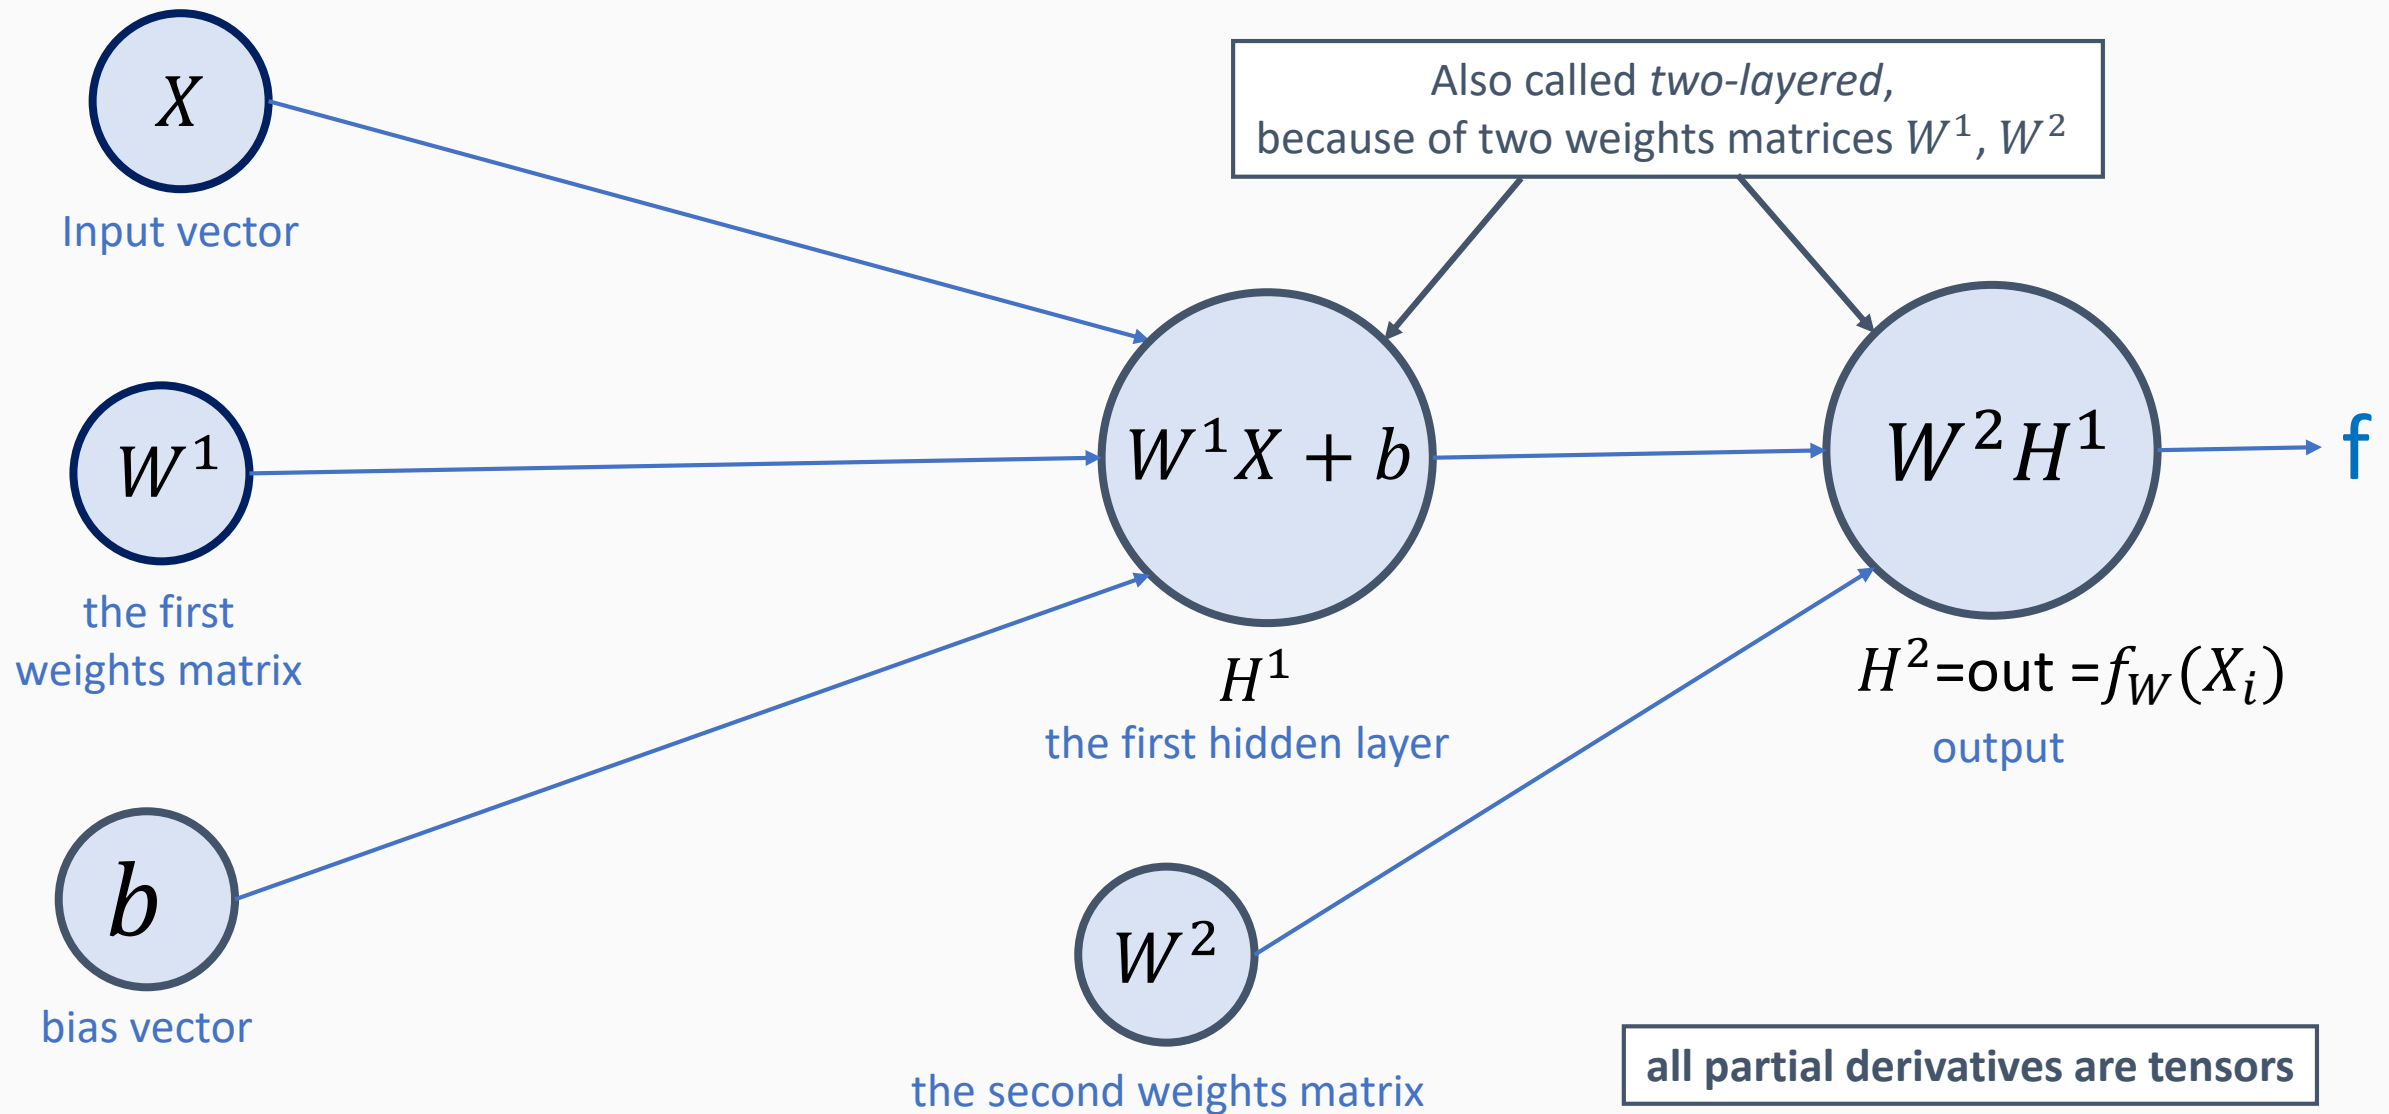
W_{jk}^i weight matrix components,

H^i matrix of hidden layer values,

multi-index

T_j a tensor (partial derivatives).

One hidden-layer neural net with linear activation



Simple network dimensions

There are t , n dimensional examples, hence $X \in \mathbb{R}^{n \times t}$, and $X_i \in \mathbb{R}^n$ is a single example

NN has n input, m hidden layer, k output neurons (denoted $NN(n, m, k)$)

Hence,

- $Y \in \mathbb{R}^{k \times t}$, $Y_i \in \mathbb{R}^k$
- $W^1 \in \mathbb{R}^{m \times n}$,
- $W^2 \in \mathbb{R}^{k \times m}$.

By convention, W is the cartesian product of weight matrices

$$W = W^1 \times W^2$$

Goal of supervised learning

Minimize

$$L(W) = \frac{1}{t} \sum_{i=1}^t \|f_W(X_i) - Y_i\|_2^2$$

take sum of squares norm, because
 X_i, Y_i and NN output are vectors.

Apply gradient descent for finding $W^* = \operatorname{argmin}\{L(W)\}$

$$W := W - \alpha \nabla L(W)$$

Compute the gradient of $L(W)$ with respect to weights.

$$\nabla L(W) = \begin{bmatrix} \nabla_{W^1} L(W), \\ \nabla_{W^2} L(W) \end{bmatrix}$$

$$W^1 := W^1 - \alpha \nabla_{W^1} L(W), \quad W^2 := W^2 - \alpha \nabla_{W^2} L(W).$$

Compute the gradient of the loss function

$$L(W) = \frac{1}{t} \sum_{i=1}^t \|f_W(X_i) - Y_i\|_2^2$$

Iterate over whole
training set $\{X_i, Y_i\}_{i=1}^t$

matrix $\nabla_{W^1} L(W) = \frac{2}{t} \sum_{i=1}^t (f_W(X_i) - Y_i) \cdot \nabla_{W^1} f_W(X_i),$

vector $\nabla_{W^2} L(W) = \frac{2}{t} \sum_{i=1}^t (f_W(X_i) - Y_i) \cdot \nabla_{W^2} f_W(X_i)$

derivative of NN with respect to weights W^1 (tensor)

derivative of NN with respect to weights W^2 (tensor)

Vector - tensor product

$$(f_W(X_i) - Y_i) \cdot \nabla_{W^1} f_W(X_i) = \sum_{l=1}^k (f_W(X_i) - Y_i)_l \nabla_{W^1} (f_W(X_i))_l$$

Tensor partial derivatives

Question: how to compute derivatives with respect to a matrix ???

$$H^2 = \text{out} = f_W(X_i) = W^2 \cdot H^1$$

$$\frac{\partial H^2}{\partial W^2} \leftarrow H^2 \text{ is a vector}$$

Derivative w.r.t. matrix W^2
(it is a tensor)

$$H^2 \in \mathbb{R}^k,$$

$$W^2 \in \mathbb{R}^{k \times m},$$

$$\frac{\partial H^2}{\partial W^2} \in \mathbb{R}^{k \times k \times m}, \leftarrow \text{a tensor}$$

$$\frac{\partial H^2}{\partial W^1} \in \mathbb{R}^{k \times m \times n}$$

Partial derivatives are tensors

We wanna compute the partial derivatives with respect to weights

$$\frac{\partial H^1}{\partial W^1}, \quad \frac{\partial H^2}{\partial W^2}, \quad \frac{\partial H^2}{\partial W^1}$$

$$H^1 = W^1 \cdot X_i$$

$$H^2 = W^2 \cdot H^1$$

Hence all of these are 3D tensors, as opposed to 2D matrix

First, we use the explicit formula for

$$\frac{\partial H_j^2}{\partial W^2} = \begin{bmatrix} \overbrace{0 \dots 0}^{m \text{ zeroes}} \\ \dots \\ 0 \dots 0 \\ (H^1)^T \\ 0 \dots 0 \\ \dots \\ 0 \dots 0 \end{bmatrix} \xleftarrow{j\text{-th row}} \begin{bmatrix} \overbrace{0 \dots 0}^{n \text{ zeroes}} \\ \dots \\ 0 \dots 0 \\ X_i^T \\ 0 \dots 0 \\ \dots \\ 0 \dots 0 \end{bmatrix} = \frac{\partial H_j^1}{\partial W^1}$$

Chain rule for tensors

How to compute the partial derivatives with respect to weights


The diagram shows the chain rule for the partial derivative of H_j^2 with respect to W^1 . The equation is
$$\frac{\partial H_j^2}{\partial W^1} = \sum_k \frac{\partial H_j^2}{\partial H_k^1} \frac{\partial H_k^1}{\partial W^1}$$
 Annotations include:

- An arrow from $\frac{\partial H_j^2}{\partial W^1}$ to the text "a matrix".
- An arrow from the summation index k to the text "Sum over all indices".
- An arrow from $\frac{\partial H_j^2}{\partial H_k^1}$ to the text "Scalar".
- An arrow from $\frac{\partial H_k^1}{\partial W^1}$ to the text "matrix".
- An arrow from the text " $j - th$ element of vector" to the j in H_j^2 .
- An arrow from the text "Scalar" to the k in H_k^1 .

 The term $\frac{\partial H_k^1}{\partial W^1}$ is circled in blue.


Repeat computation for all indices j of H^2

Backprop compute $\nabla_{W^1} f_W(X_i), \nabla_{W^2} f_W(X_i)$




Verify your backprop implementation using gradient checking

Given X_i



Backprop to compute the partial derivatives

Estimate the rate of change of $f_W(X_i)$
for fixed (small) h in all W^1, W^2 directions



\approx
comparable

$$\frac{\partial f_W(X_i)}{\partial W} \approx \frac{f_{W+h}(X_i) - f_W(X_i)}{h},$$

where $W + h$ means we add h
to a single component of W
(and repeat it for all i to compute
the whole partial derivatives tensor)

Obtaining tensor partial *numerical* derivative (for gradient checking)

For example to compute numerical derivative $\approx \frac{\partial f_W(X_i)}{\partial W^1}$
need to estimate the finite difference in all possible directions,
i.e. for all $1 \leq i \leq m, 1 \leq j \leq n$ compute modified weights W^1 by

$W^{1*} = W^1$, where for the selected entry we set $W_{ij}^{1*} = W_{ij}^1 + h$, we obtain
whole slice of the partial derivatives tensor,
remembering $W = W^1 \times W^2$, we set $W^* = W^{1*} \times W^2$


$$\frac{\partial f_W(X_i)}{\partial W_{ij}^1} \approx \frac{f_{W^*}(X_i) - f_W(X_i)}{h}$$

And repeat this for all i, j

Weight matrices pair having
 i, j -th entry of W^1 modified

Value of this is a vector
[:,i,j] slice of the partial
derivative tensor

$\frac{\partial f_W(X_i)}{\partial W^1}$ 3D tensor




element of the tensor indexed by
(k, i, j) is


$$\frac{\partial f_W(X_i)_k}{\partial W_{ij}^1}$$

i.e. dimension k is the dimension
of vector, i, j are dimensions of
weight matrices.

Concrete example $\nabla_{W^1} f_W(X_i), \nabla_{W^2} f_W(X_i), b = 0$




Edge $W^2 - H^2$



In this case $\frac{\partial H^2}{\partial W^2}$ is a $3 \times 3 \times 3$ dimensional tensor

$$\frac{\partial H_1^2}{\partial W^2} = \begin{bmatrix} 6 & 15 & 24 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix} \quad \frac{\partial H_2^2}{\partial W^2} = \begin{bmatrix} 0 & 0 & 0 \\ 6 & 15 & 24 \\ 0 & 0 & 0 \end{bmatrix} \quad \frac{\partial H_3^2}{\partial W^2} = \begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & 0 \\ 6 & 15 & 24 \end{bmatrix}$$

$$\frac{\partial H^2}{\partial W^2} = \begin{bmatrix} 6 & 15 & 24 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix} \begin{bmatrix} 0 & 0 & 0 \\ 6 & 15 & 24 \\ 0 & 0 & 0 \end{bmatrix} \begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & 0 \\ 6 & 15 & 24 \end{bmatrix}$$



Other components of the tensor

$$\frac{\partial H_2^2}{\partial W^1} = \sum_k \frac{\partial H_2^2}{\partial H_k^1} \frac{\partial H_k^1}{\partial W^1} = 0 \cdot \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix} + 1 \cdot \begin{bmatrix} 0 & 0 & 0 \\ 1 & 1 & 1 \\ 0 & 0 & 0 \end{bmatrix} + 1 \cdot \begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & 0 \\ 1 & 1 & 1 \end{bmatrix}$$

$$\frac{\partial H_3^2}{\partial W^1} = \sum_k \frac{\partial H_3^2}{\partial H_k^1} \frac{\partial H_k^1}{\partial W^1} = 1 \cdot \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix} + 0 \cdot \begin{bmatrix} 0 & 0 & 0 \\ 1 & 1 & 1 \\ 0 & 0 & 0 \end{bmatrix} + 1 \cdot \begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & 0 \\ 1 & 1 & 1 \end{bmatrix}$$

$$\frac{\partial H^2}{\partial W^1} = \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 0 & 0 & 0 \end{bmatrix} \begin{bmatrix} 0 & 0 & 0 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix}$$

Compute the gradient of the loss function

$$L(W) = \frac{1}{t} \sum_{i=1}^t \|f_W(X_i) - Y_i\|_2^2$$

By chain rule

$$\nabla_{W^1} L(W) = \frac{2}{t} \sum_{i=1}^t (f_W(X_i) - Y_i) \cdot \nabla_{W^1} f_W(X_i),$$

derivative of NN with
respect to weights
W1 (tensor)
(inner function)
backprop

$$\nabla_{W^2} L(W) = \frac{2}{t} \sum_{i=1}^t (f_W(X_i) - Y_i) \cdot \nabla_{W^2} f_W(X_i)$$

derivative of NN
with respect to
weights W2 (tensor)
(inner function)
backprop

Vector - tensor product

$$(f_W(X_i) - Y_i) \cdot \nabla_{W^1} f_W(X_i) = \sum_{l=1}^k (f_W(X_i) - Y_i)_l \nabla_{W^1} (f_W(X_i))_l$$

Gradient of the loss in practice

$$(f_W(X_i) - Y_i) \cdot \nabla_{W^1} f_W(X_i) = \sum_{l=1}^k (f_W(X_i) - Y_i)_l \nabla_{W^1} (f_W(X_i))_l$$

Assume training pair

$$(X_i, Y_i) = \left(\begin{bmatrix} 1 \\ 1 \\ 1 \end{bmatrix}, \begin{bmatrix} 20 \\ 0 \\ 0 \end{bmatrix} \right) \text{ and } f_W(X_i) = \begin{bmatrix} 21 \\ 39 \\ 30 \end{bmatrix}$$

Recall W^1 gradient of the output

$$\frac{\partial H^2}{\partial W^1} = \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 0 & 0 & 0 \end{bmatrix} \begin{bmatrix} 0 & 0 & 0 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ 0 & 1 & 1 \end{bmatrix}$$

Then, in this case (for the given training pair (X_i, Y_i))

$$(f_W(X_i) - Y_i) \cdot \nabla_{W^1} f_W(X_i) = 1 \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \\ 0 & 0 & 0 \end{bmatrix} + 39 \begin{bmatrix} 0 & 0 & 0 \\ 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} + 30 \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ 1 & 1 & 1 \end{bmatrix}$$

We know everything to perform NN learning

$$(f_W(X_i) - Y_i) \cdot \nabla_{W^1} f_W(X_i) = \begin{bmatrix} 31 & 31 & 31 \\ 40 & 40 & 40 \\ 69 & 69 & 69 \end{bmatrix}$$

Repeat for all pairs in the training set and sum up,
this is the final gradient $\nabla_{W^1} L(W)$.

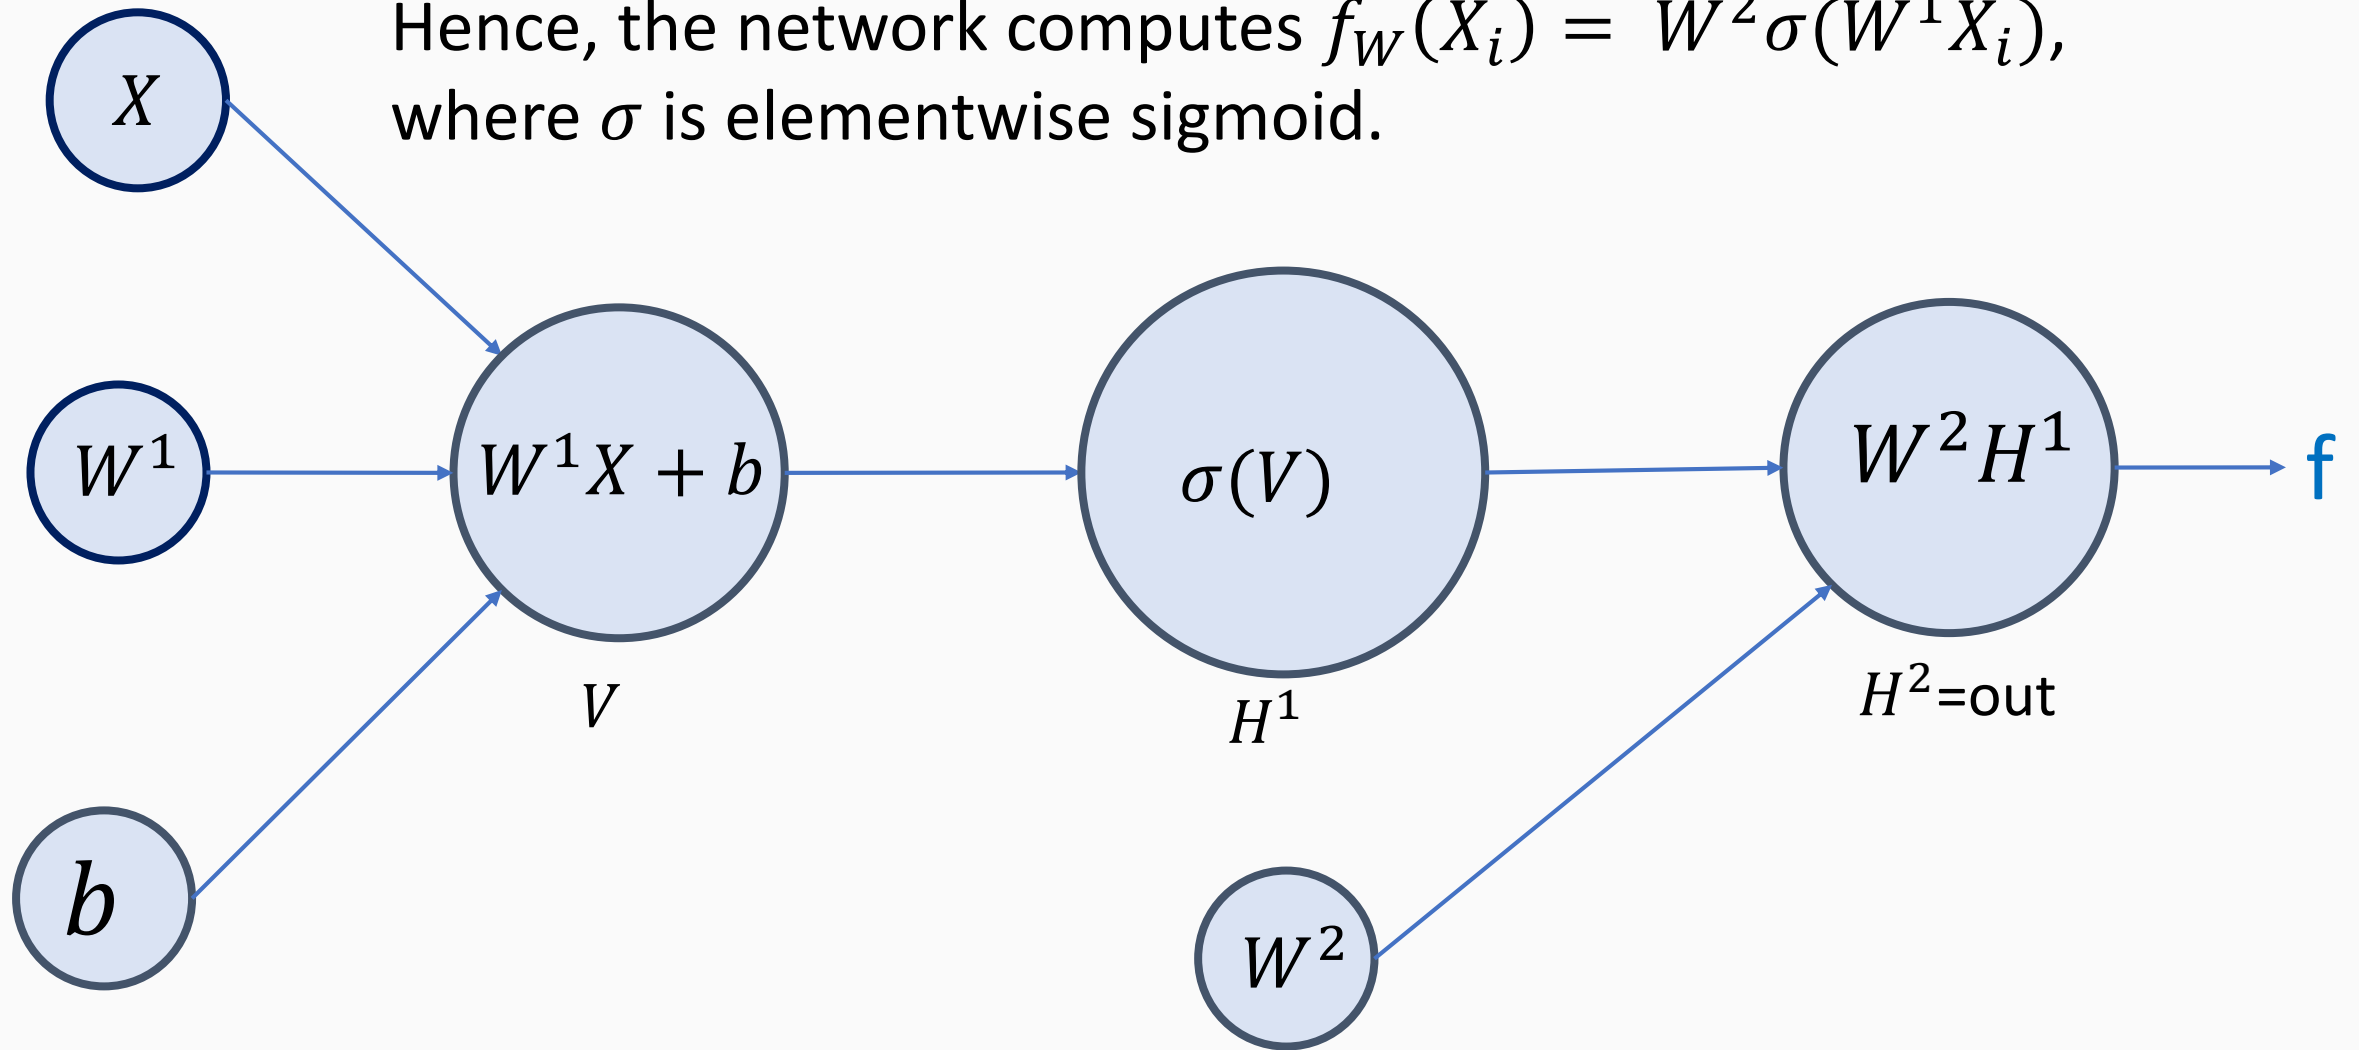
$$\nabla_{W^1} L(W) = \frac{2}{t} \sum_{i=1}^t (f_W(X_i) - Y_i) \cdot \nabla_{W^1} f_W(X_i)$$

Compute $\nabla_{W^2} L(W)$ the same way.
And perform the gradient descent update

$$W^1 := W^1 - \alpha \nabla_{W^1} L(W), \quad W^2 := W^2 - \alpha \nabla_{W^2} L(W).$$

Now the hidden layer computes sigmoids

Hence, the network computes $f_W(X_i) = W^2 \sigma(W^1 X_i)$, where σ is elementwise sigmoid.



The same loss, but sigmoidal network

$$L(W) = \frac{1}{t} \sum_{i=1}^t \|f_W(X_i) - Y_i\|_2^2$$

$$\nabla_{W^1} L(W) = \frac{2}{t} \sum_{i=1}^t (f_W(X_i) - Y_i) \cdot \nabla_{W^1} f_W(X_i)$$

Gradient of sigmoidal NN

$$f_W(X_i) = W^2 (\sigma(W^1 X_i + b))$$

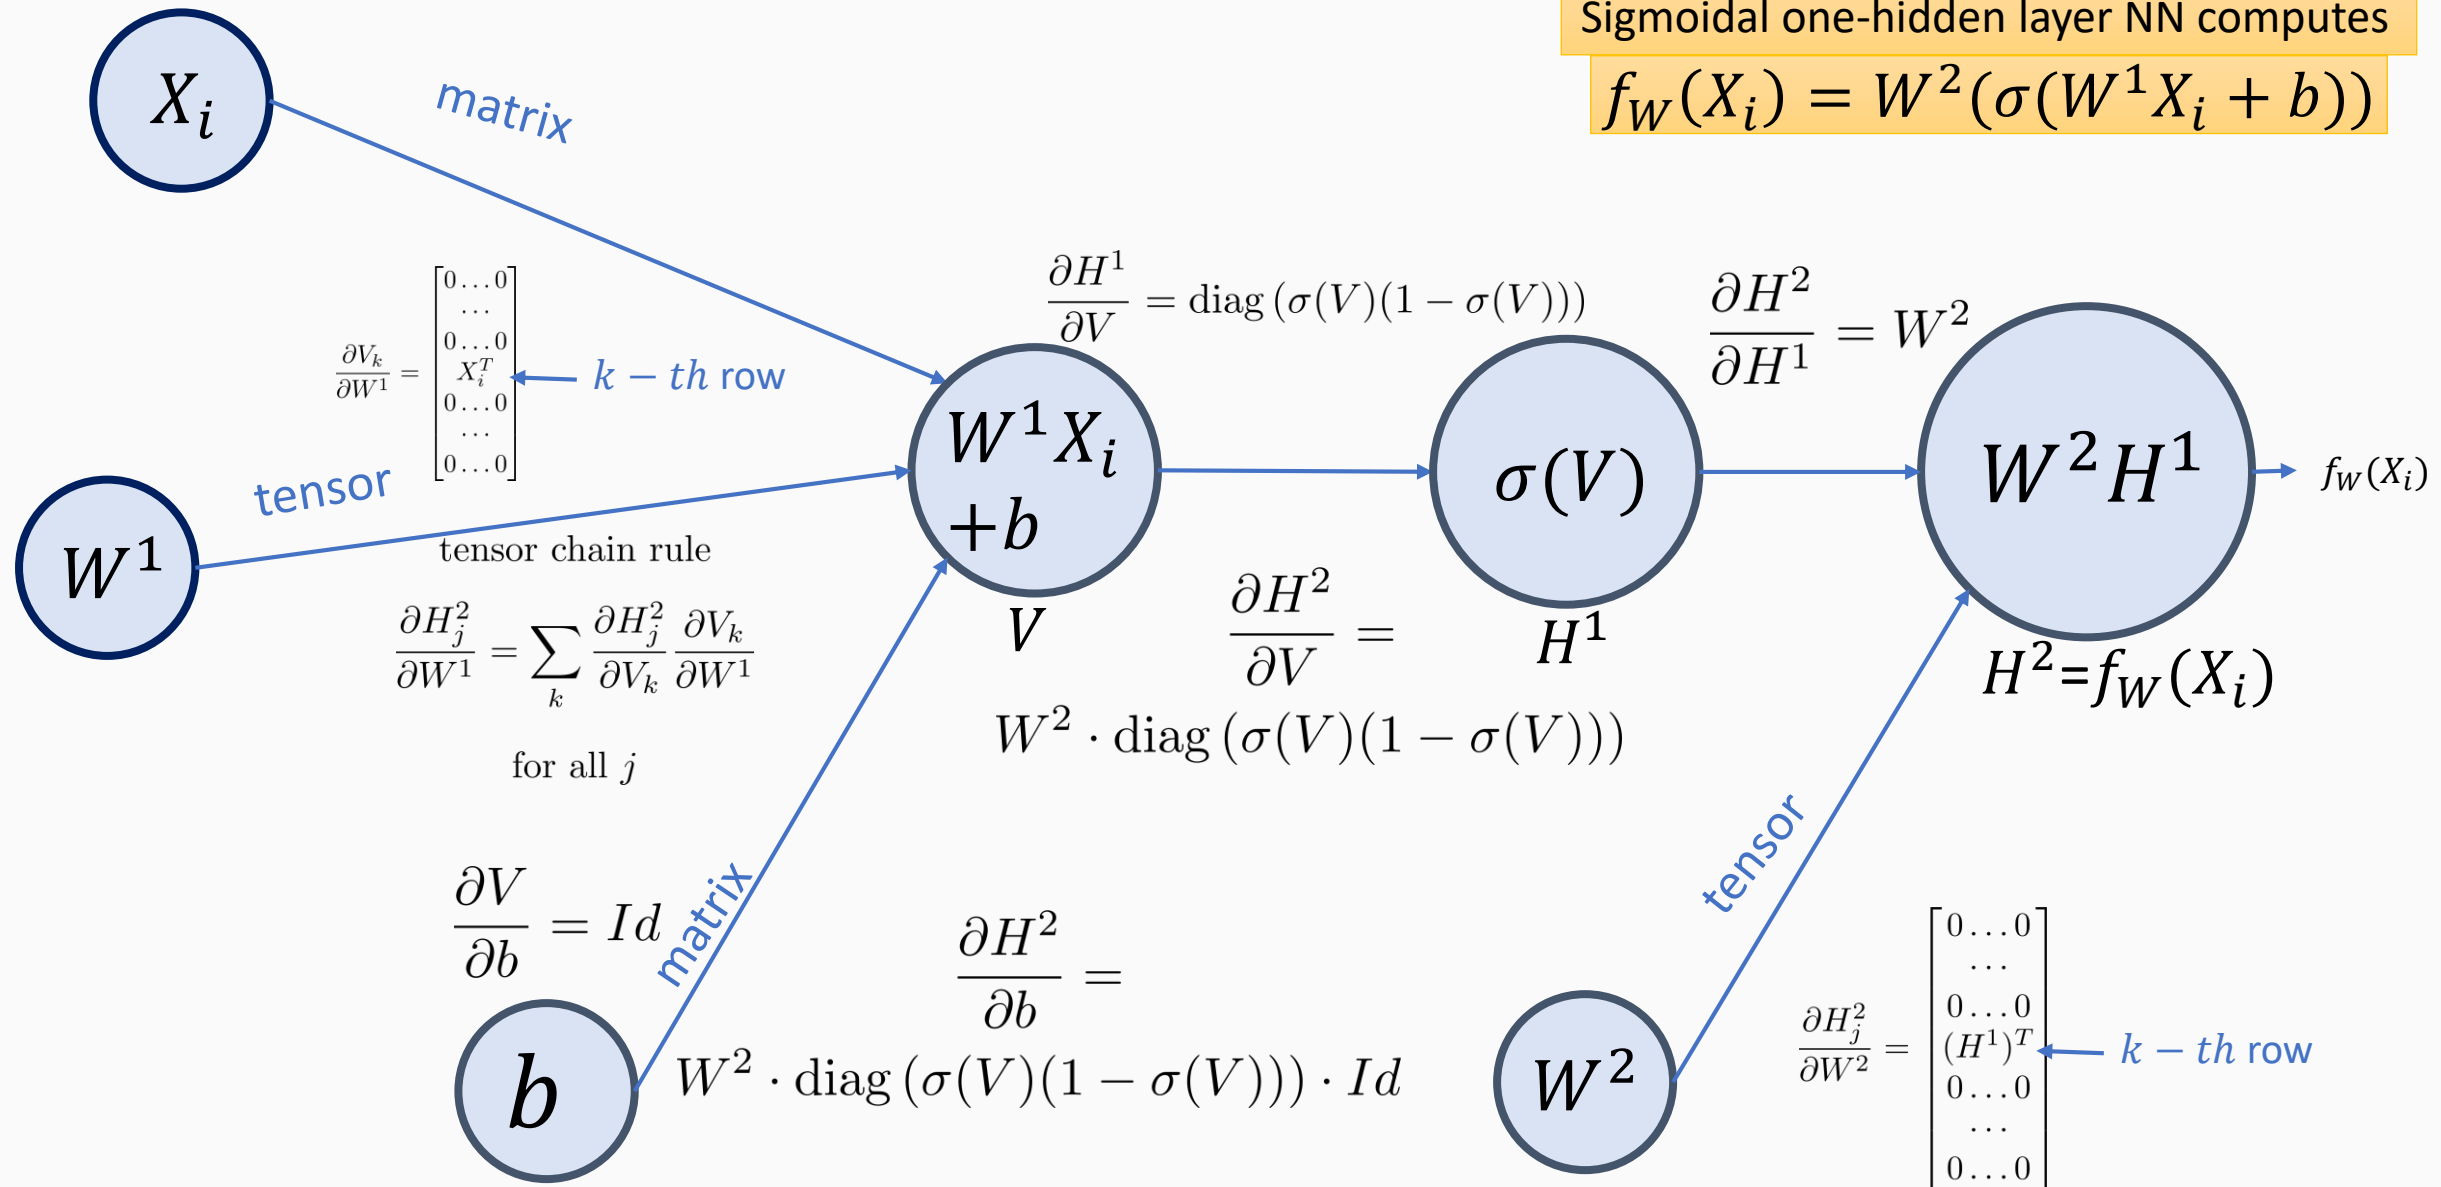
Element-wise sigmoid function

The next slides present how to compute using backprop gradients of sigmoidal net $\nabla_{W^1} f_W(X_i)$, $\nabla_{W^2} f_W(X_i)$

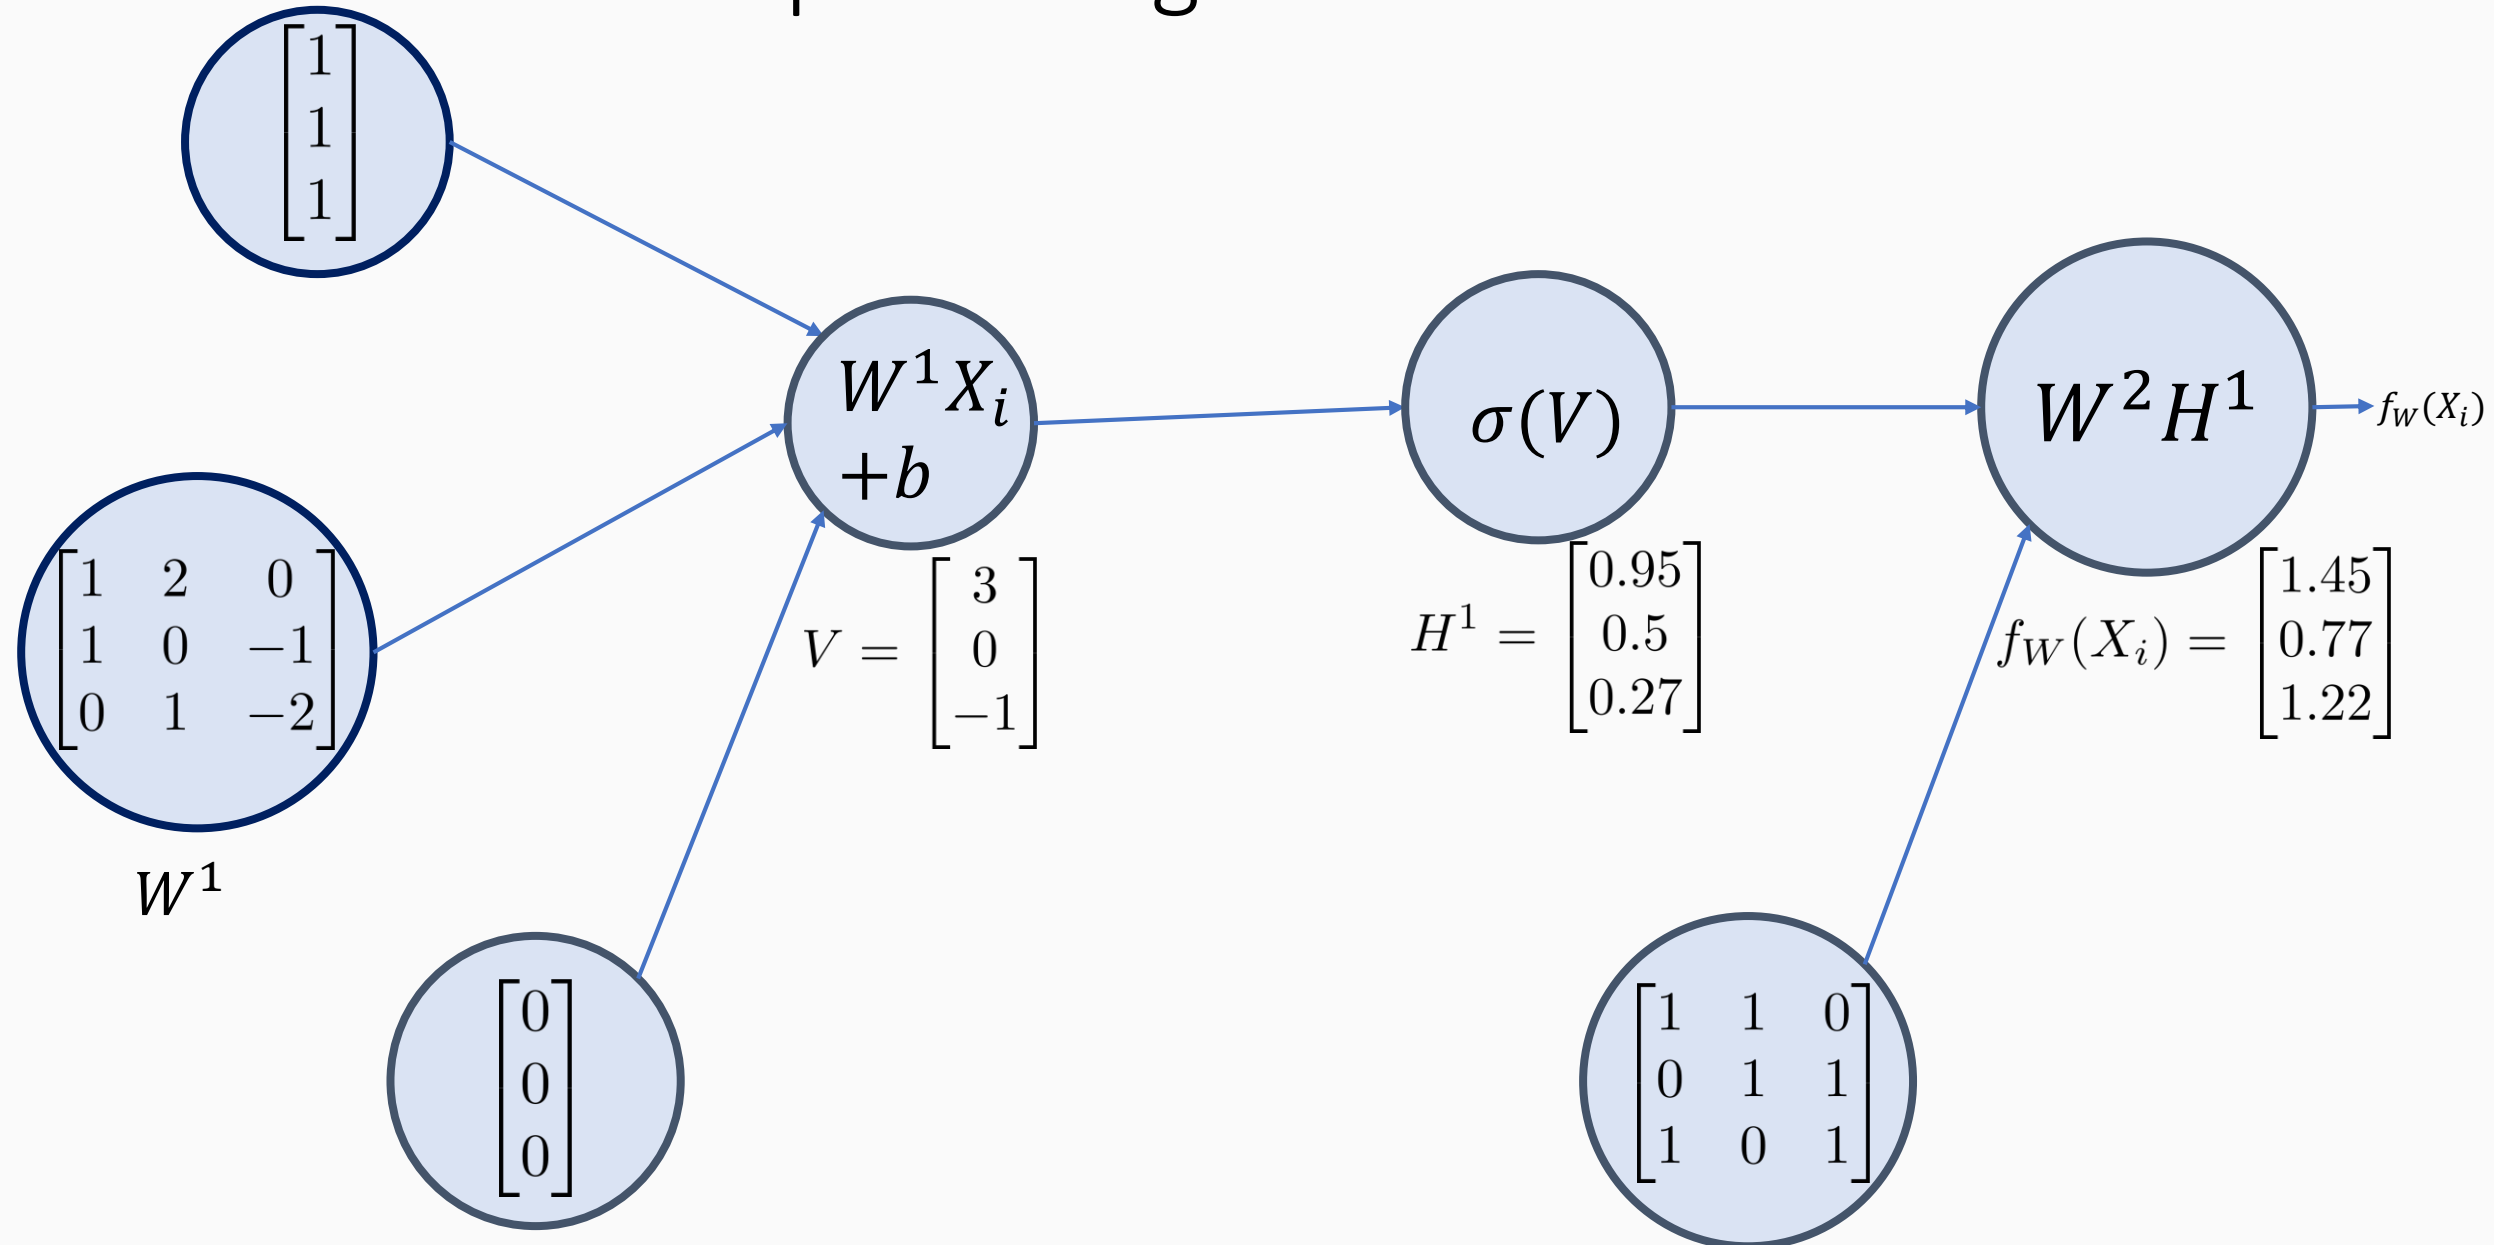
Backprop compute $\nabla_{W^1} f_W(X_i)$, $\nabla_{W^2} f_W(X_i)$ for sigmoidal net

Sigmoidal one-hidden layer NN computes

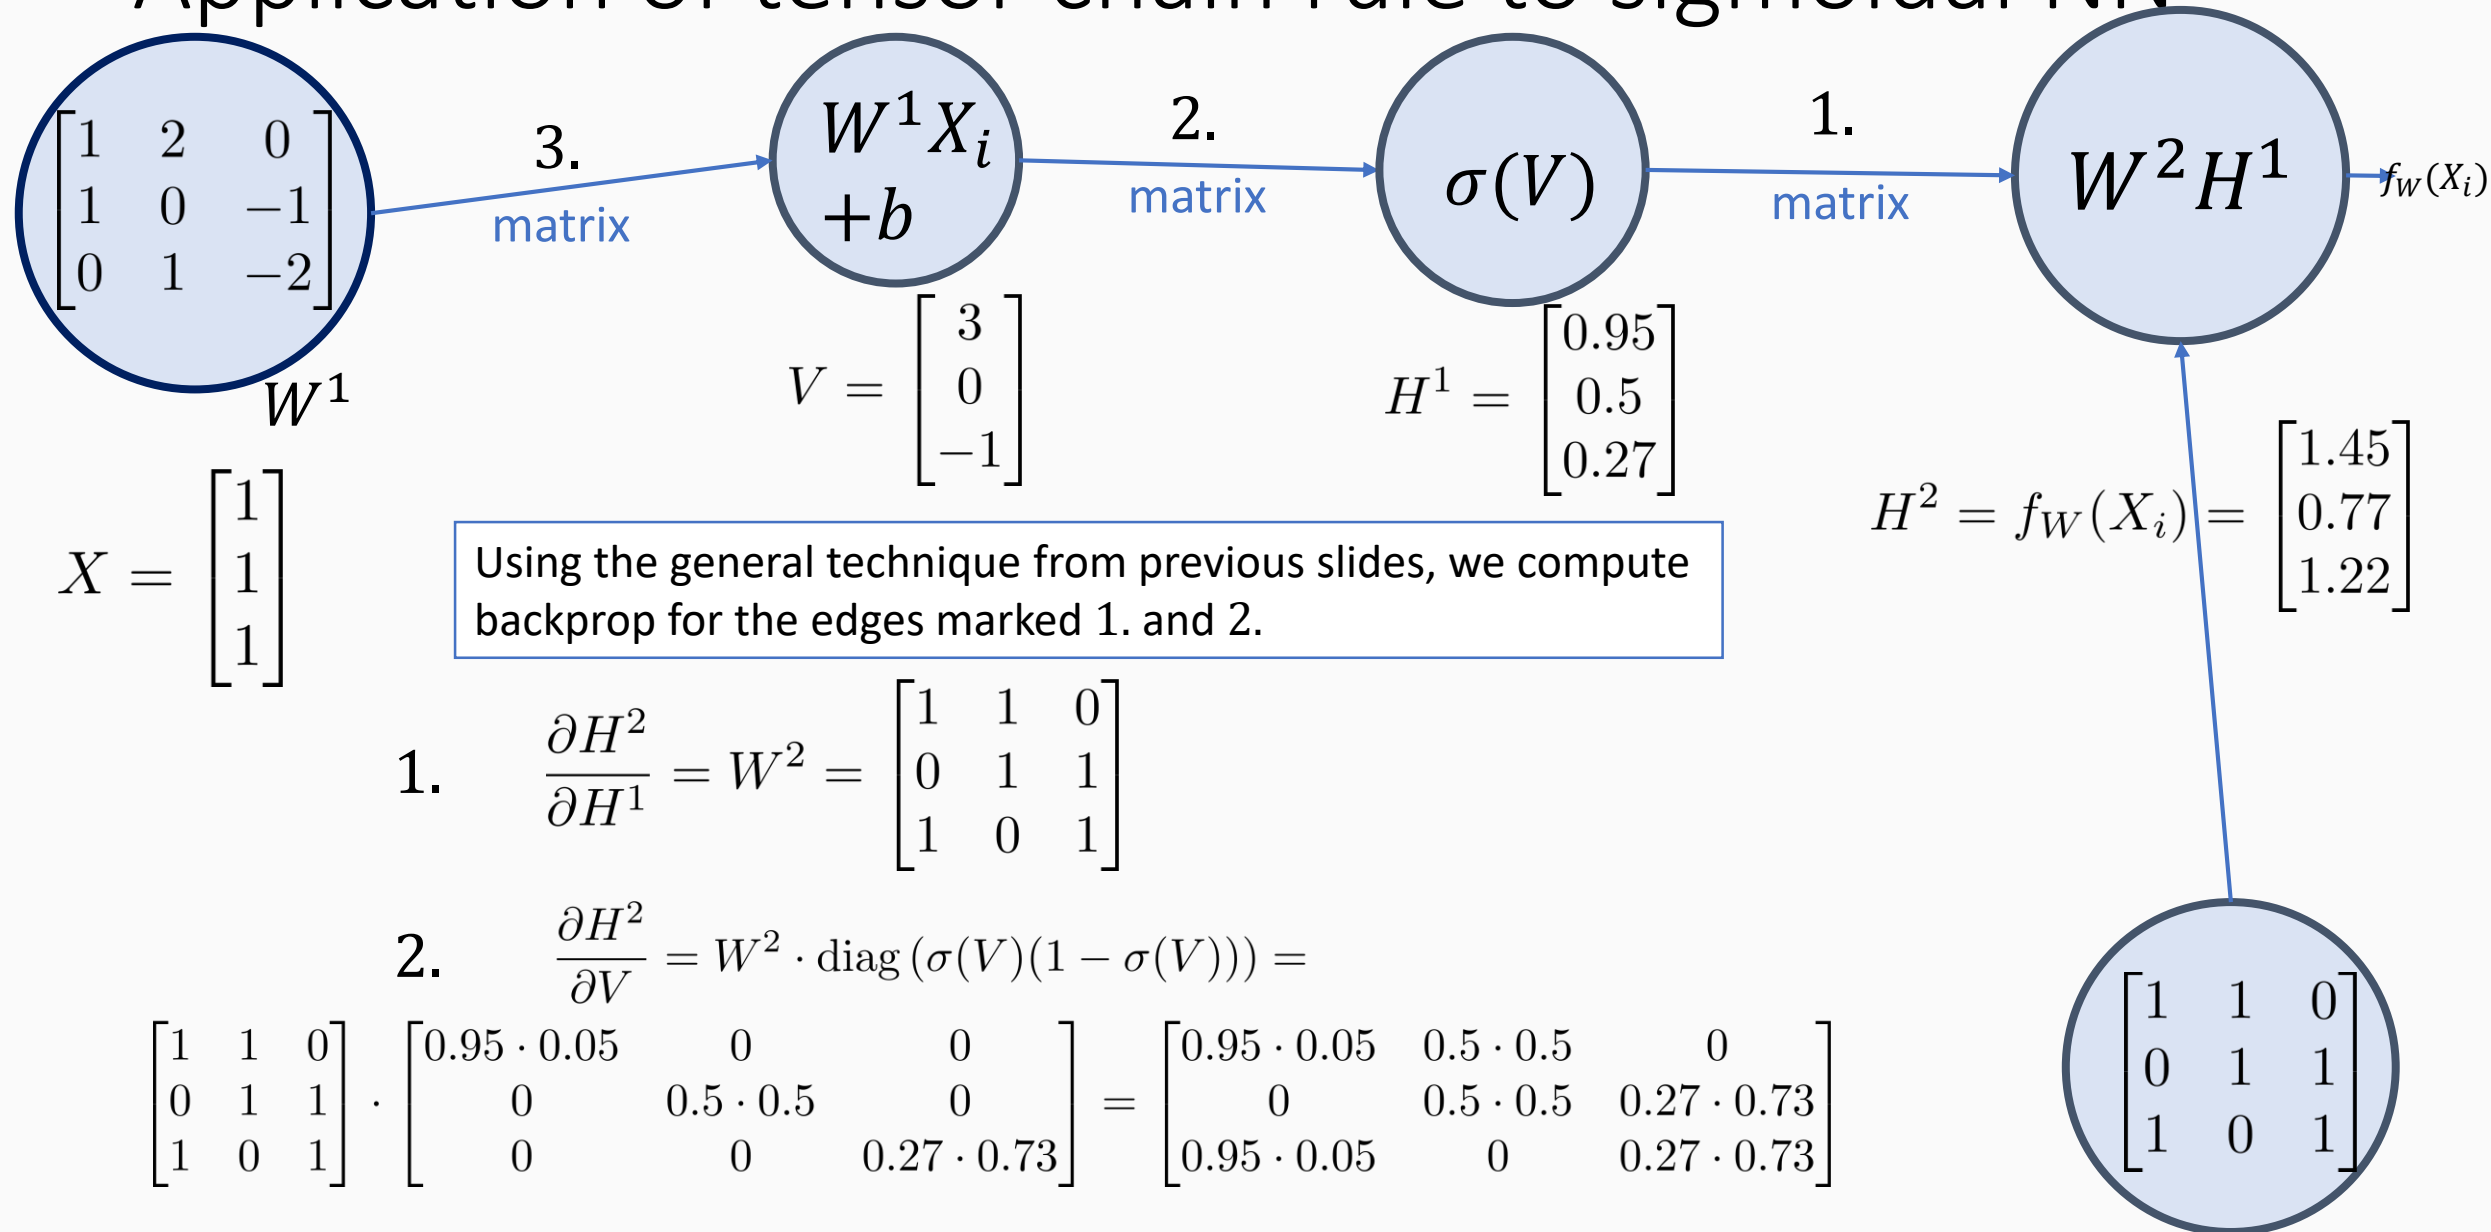
$$f_W(X_i) = W^2(\sigma(W^1 X_i + b))$$



Concrete example for sigmoidal network



Application of tensor chain rule to sigmoidal NN



Compute the tensor partial derivative for edge 3.

$$H^2 = f_W(X_i)$$

We apply the general formulas

From step 2. we have

3.

tensor chain rule

$$\frac{\partial V_k}{\partial W^1} = \begin{bmatrix} 0 \dots 0 \\ \vdots \\ 0 \dots 0 \\ X_i^T \\ 0 \dots 0 \\ \vdots \\ 0 \dots 0 \end{bmatrix} \quad \leftarrow k - th \text{ row}$$

$$\frac{\partial H_j^2}{\partial W^1} = \sum_k \frac{\partial H_j^2}{\partial V_k} \frac{\partial V_k}{\partial W^1} \quad \text{for all } j$$

$$\frac{\partial H^2}{\partial V} = \begin{bmatrix} 0.95 \cdot 0.05 & 0.5 \cdot 0.5 & 0 \\ 0 & 0.5 \cdot 0.5 & 0.27 \cdot 0.73 \\ 0.95 \cdot 0.05 & 0 & 0.27 \cdot 0.73 \end{bmatrix}$$

$$\frac{\partial H_1^2}{\partial W^1} = 0.95 \cdot 0.05 \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix} + 0.5 \cdot 0.5 \begin{bmatrix} 0 & 0 & 0 \\ 1 & 1 & 1 \\ 0 & 0 & 0 \end{bmatrix} = \begin{bmatrix} 0.95 \cdot 0.05 & 0.95 \cdot 0.05 & 0.95 \cdot 0.5 \cdot 0.5 \\ 0.5 \cdot 0.5 & 0.5 \cdot 0.5 & 0.5 \cdot 0.5 \\ 0 & 0 & 0 \end{bmatrix},$$

$$\frac{\partial H_2^2}{\partial W^1} = 0.5 \cdot 0.5 \begin{bmatrix} 0 & 0 & 0 \\ 1 & 1 & 1 \\ 0 & 0 & 0 \end{bmatrix} + 0.27 \cdot 0.73 \begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & 0 \\ 1 & 1 & 1 \end{bmatrix} = \begin{bmatrix} 0 & 0 & 0 \\ 0.5 \cdot 0.5 & 0.5 \cdot 0.5 & 0.5 \cdot 0.5 \\ 0.27 \cdot 0.73 & 0.27 \cdot 0.73 & 0.27 \cdot 0.73 \end{bmatrix},$$

$$\frac{\partial H_3^2}{\partial W^1} = 0.95 \cdot 0.05 \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ 0 & 0 & 0 \end{bmatrix} + 0.27 \cdot 0.73 \begin{bmatrix} 0 & 0 & 0 \\ 0 & 0 & 0 \\ 1 & 1 & 1 \end{bmatrix} = \begin{bmatrix} 0.95 \cdot 0.05 & 0.95 \cdot 0.05 & 0.95 \cdot 0.05 \\ 0 & 0 & 0 \\ 0.27 \cdot 0.73 & 0.27 \cdot 0.73 & 0.27 \cdot 0.73 \end{bmatrix}$$

Combining results


This will give $\nabla_{W^1} f_W(X_i)$, $\nabla_{W^2} f_W(X_i)$ for sigmoidal net.

To compute the gradients of the loss function still need to perform the computation presented on slides **39, 40, 41** (combining derivative of the loss with the gradient of NN).

ReLU net

$$\text{ReLU}(x) = \begin{cases} x & \text{if } x > 0, \\ 0 & \text{if } x \leq 0 \end{cases},$$

$$\text{ReLU}'(x) = \begin{cases} 1 & \text{if } x > 0, \\ 0 & \text{if } x \leq 0 \end{cases}$$



Rectified Linear Units Improve Restricted Boltzmann Machines

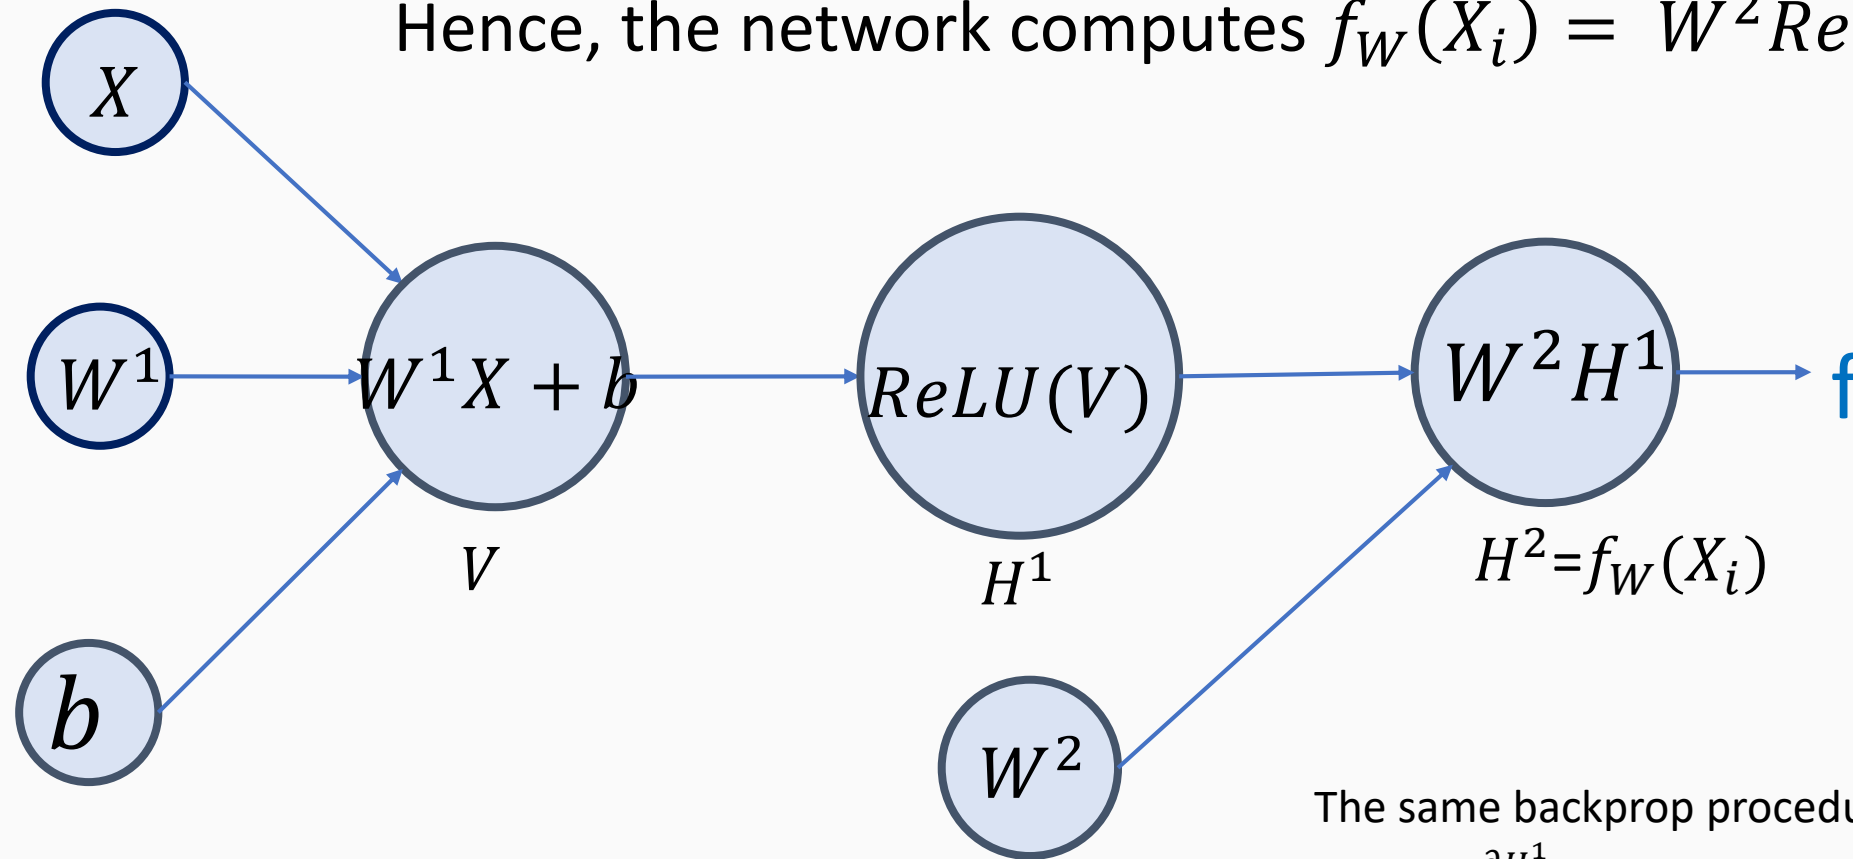
Vinod Nair
Geoffrey E. Hinton

Department of Computer Science, University of Toronto, Toronto, ON M5S 2G4, Canada

VNAIR@CS.TORONTO.EDU
HINTON@CS.TORONTO.EDU

ReLU net

Hence, the network computes $f_W(X_i) = W^2 \text{ReLU}(W^1 X_i)$



The same backprop procedure as for sigmoidal net, but $\frac{\partial H_i^1}{\partial v_i} = 1$ if $\text{ReLU}(V_i) > 0$,
0 otherwise.

Another loss – Cross entropy loss

$$L(W) = - \sum_{i=1}^t Y_i \cdot \log (f_W(X_i))$$

Desired output
(usually binary vector
having one element = 1)

Gradient of sigmoidal NN

Observe that the outputs of NN are now interpreted as probabilities ,
in particular it holds that $f_w(X_i) \in [0,1]$ (the loss stays positive),

need to normalize i.e. an additional layer computing $f_w(X_i) / ||f_w(X_i)||$,
is required, i.e. normalized output (like *softmax*), but this
requires one additional step in backprop.

Analytic formula for global minima of (linear) NN

Neural Networks and Principal Component Analysis: Learning from Examples Without Local Minima

PIERRE BALDI AND KURT HORNIK*

University of California, San Diego

(Received 18 May 1988; revised and accepted 16 August 1988)




FIGURE 2. The landscape of E .

$$A = U_{\mathcal{Y}} C$$

$$B = C^{-1} U_{\mathcal{Y}}' \Sigma_{YX} \Sigma_{XX}^{-1}$$

For such a critical point we have

$$W = P_{U_{\mathcal{Y}}} \Sigma_{YX} \Sigma_{XX}^{-1}$$

$$E(A, B) = \text{tr} \Sigma_{YY} - \sum_{i \in \mathcal{J}} \lambda_i.$$

$$\Sigma_{YX} = Y^T X,$$

$$\Sigma_{XX} = X^T X,$$

$$U = [u_1, u_2, \dots, u_p]$$


matrix of p eigenvectors, corresponding to the p largest eigenvalues of

$$\Sigma = \Sigma_{YX} \Sigma_{XX}^{-1} \Sigma_{XY}$$

Sigmoidal vs. ReLU (One layer)

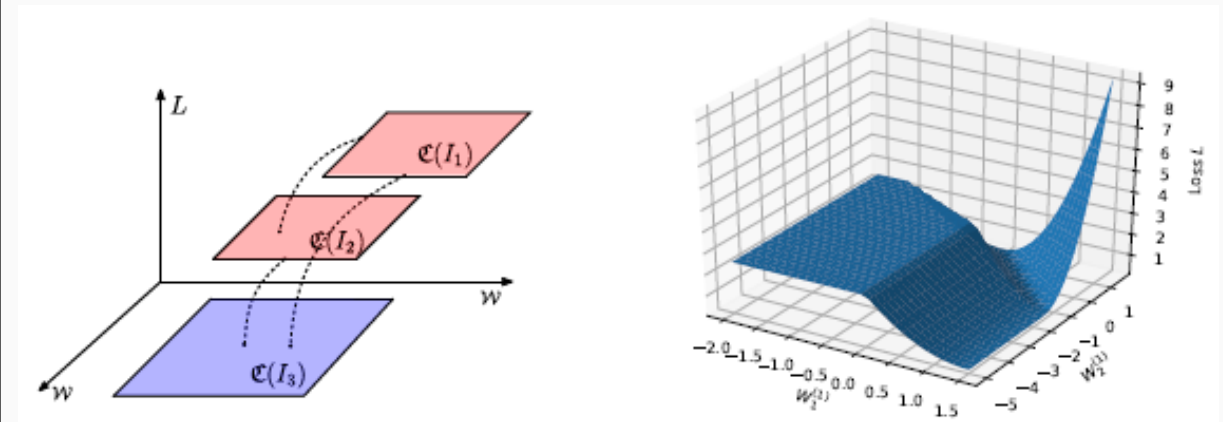
Sigmoidal

Easy to train



ReLU

Hard to train



Tensors in NumPy

```
import numpy as np
#tensors in NumPy
```

```
#how to define a tensor
```

```
n = 3
```

```
m = 3
```

```
A = np.ones( (n, n, n) ) # n,m,k dimensional tensor
```

```
B = np.ones( (n, n, n) )
```

```
#tensor product
```

```
C = np.tensordot(A, B, 0) (its like Cartesian product of two tensors)
print(C.shape) (3,3,3,3,3,3)
```

```
C = np.tensordot(A, B, 1)
print(C.shape) (3,3,3,3)
```

$$C_{i,j,l,m} = \sum_{k=1}^n A_{i,j,k} B_{k,l,m}$$

```
C = np.tensordot(A, B, 2)
print(C.shape) (3,3)
```

$$C_{i,m} = \sum_{j=1}^n \sum_{k=1}^n A_{i,j,k} B_{k,j,m}$$

```
C = np.tensordot(A, B, ([0, 1], [1, 2]))
print(C.shape) (3,3)
```

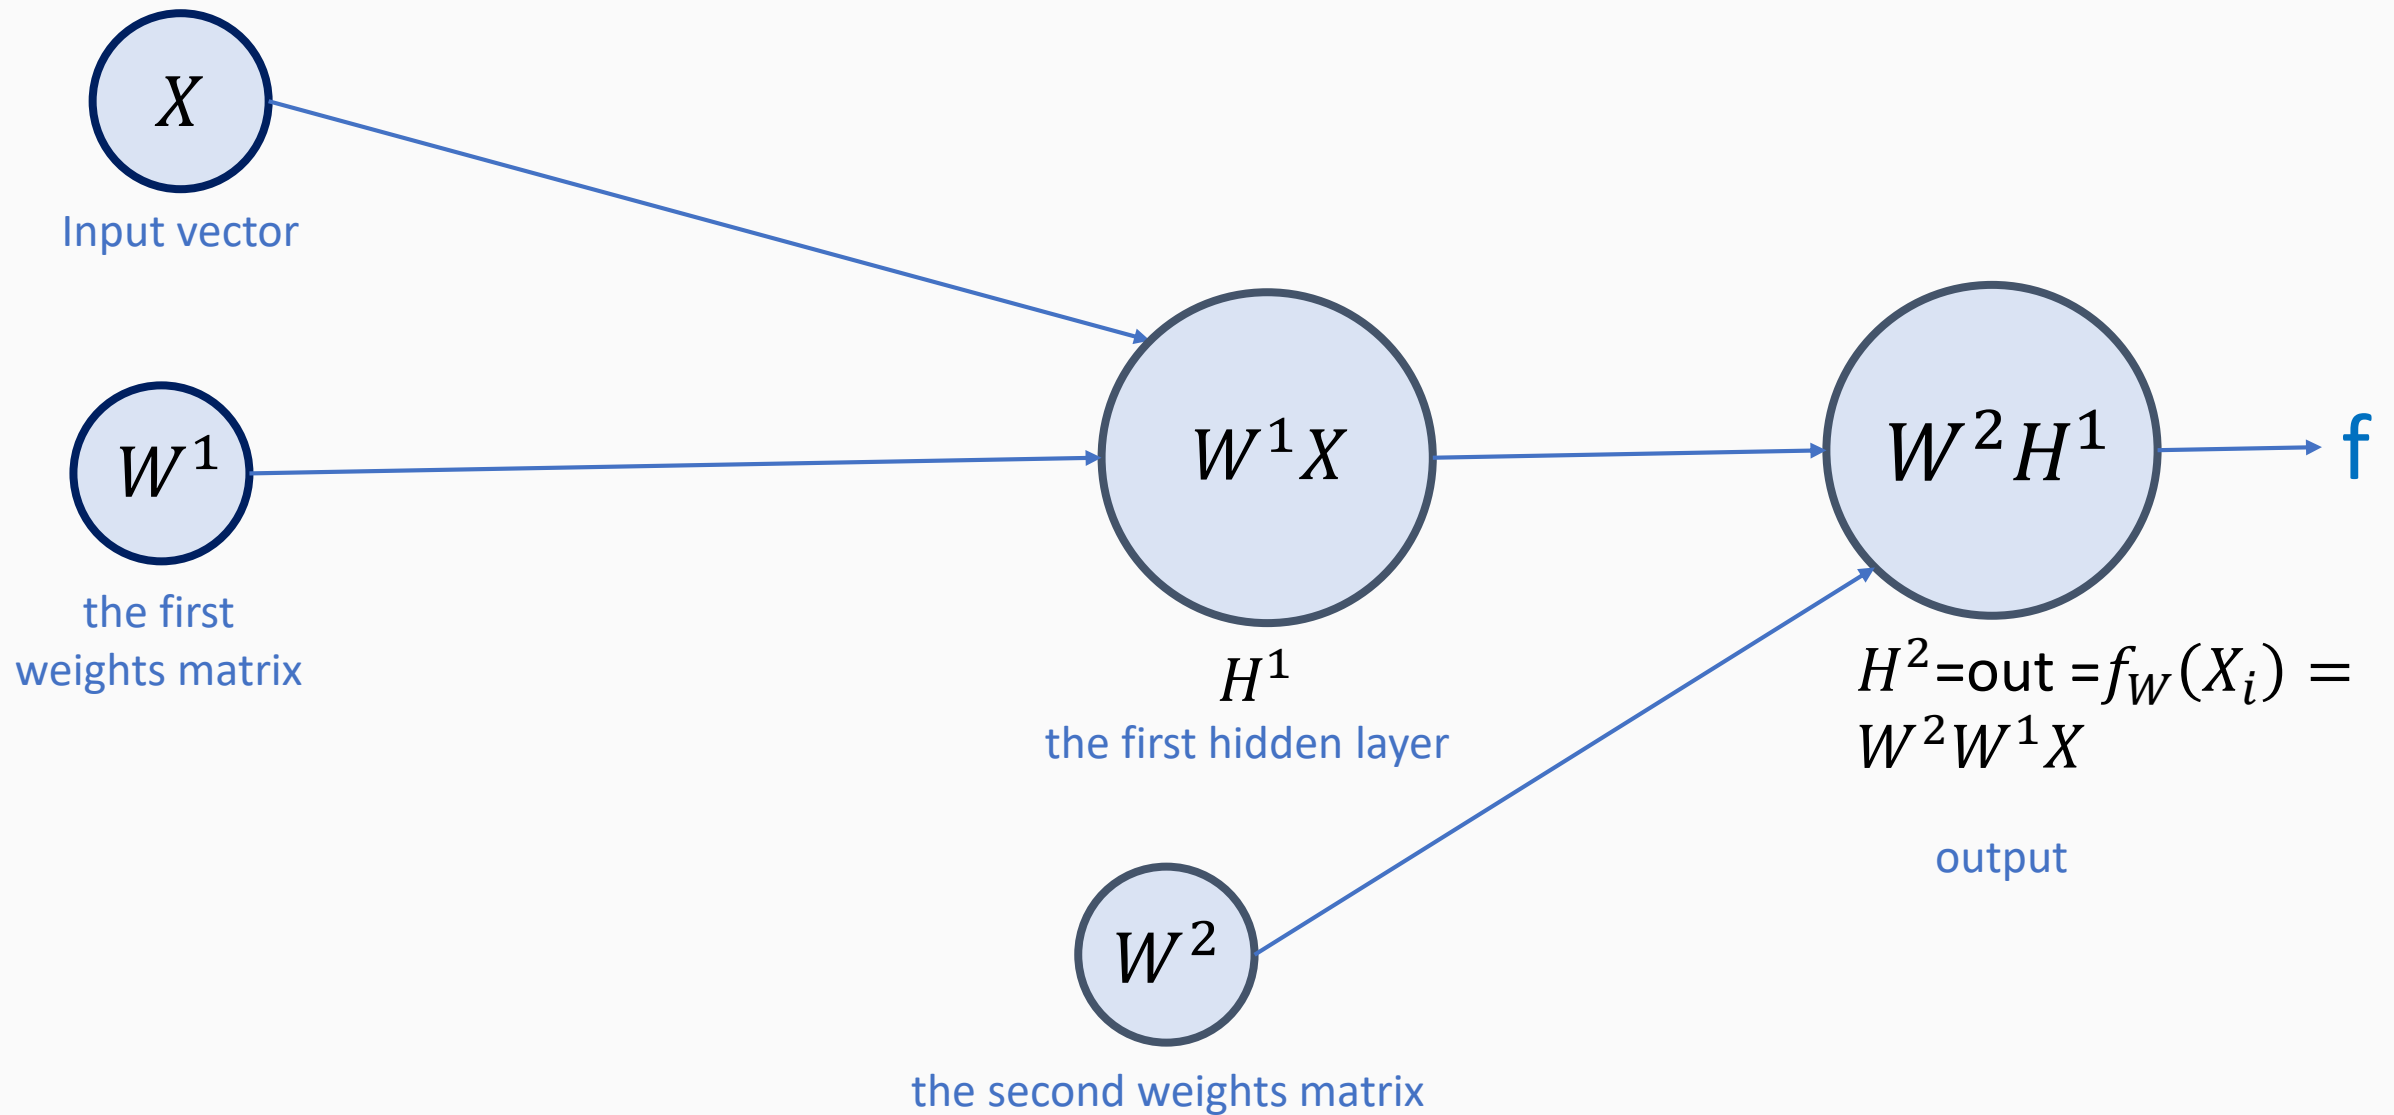
$$C_{i,m} = \sum_{j=1}^n \sum_{k=1}^n A_{j,k,i} B_{m,j,k} \text{ (any combination possible)}$$

Better to use loops to start with instead of tensordot

Remark

For coding tensor operations I suggest using just loops in python instead of tensordot function.

Assignment 8 NN1 architecture



Assignment 8 NN2 architecture

